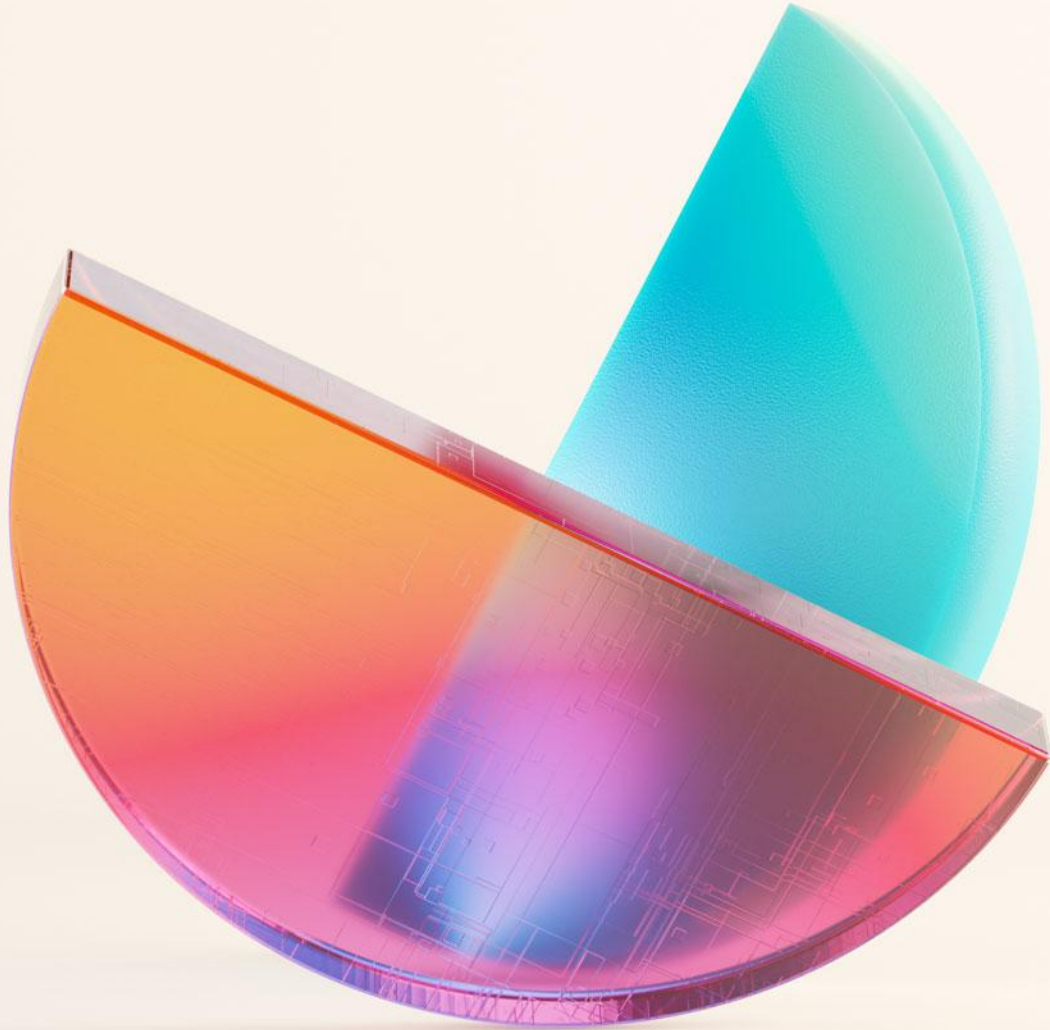




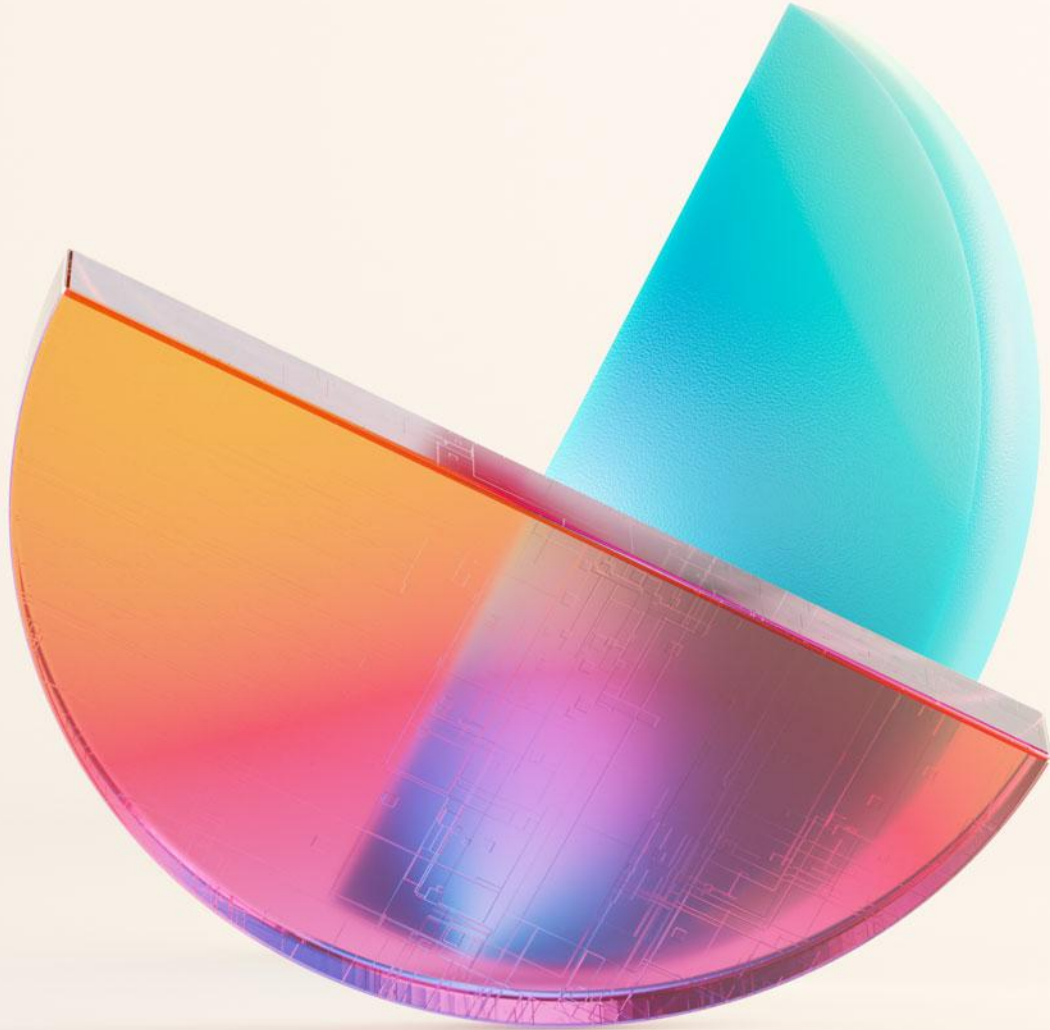
# Microsoft Fabric Community Conference



# End to End Data Engineering

Modern Data Warehousing on  
Microsoft Fabric

Wifi Network: MGMResorts-WiFi  
No password



## Join Us : Fabric User Panel



Influence our product roadmap and ensure  
Fabric meets your real-life needs

Join us at [aka.ms/FabricUserPanel](https://aka.ms/FabricUserPanel)

Microsoft Fabric  
Community Conference

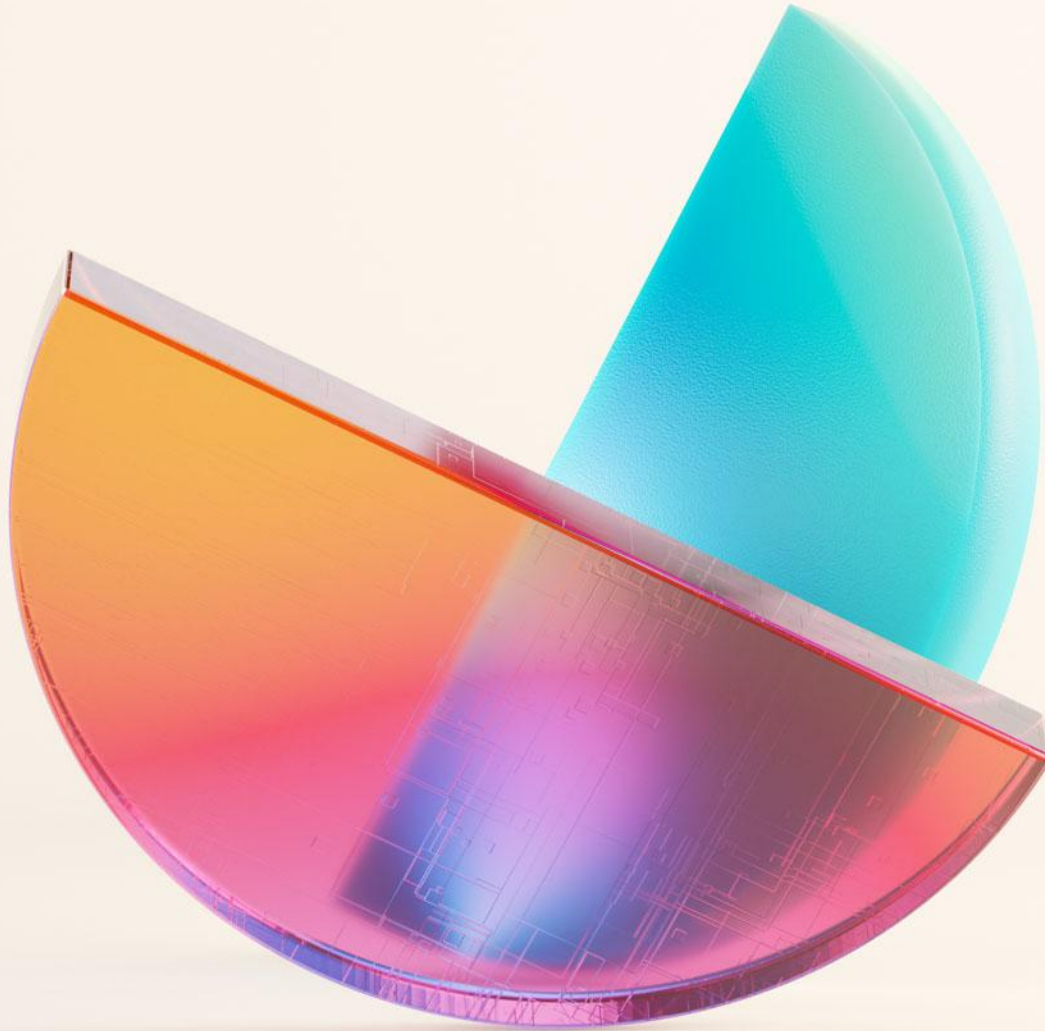
## Fast-track your career in data and AI

**Become a Microsoft Certified Fabric Data  
Engineer Associate – take Exam DP-700  
for free!**

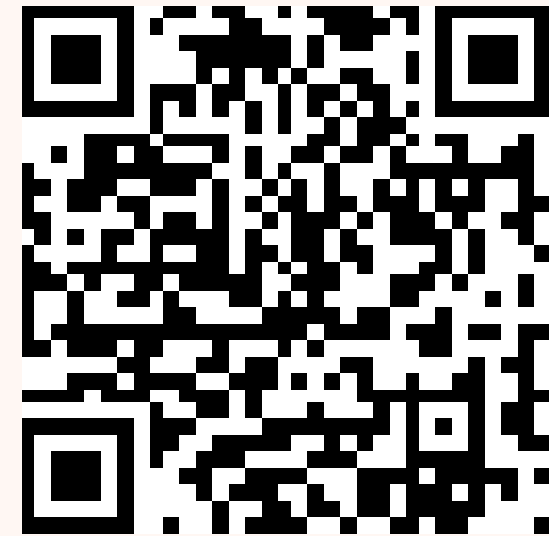
Visit the Fabric Community Lounge to learn  
more about this *limited-time offer*.



**[aka.ms/fabcon/dp700](https://aka.ms/fabcon/dp700)**



Learn more about  
Microsoft Fabric



Power your AI transformation with a  
complete data platform





# Get Involved in the Fabric Community



[aka.ms/FabricCommunity](https://aka.ms/FabricCommunity)

Connect with community members, ask questions, and learn more about Fabric



[aka.ms/FabricUserGroups](https://aka.ms/FabricUserGroups)

Find a user group that matches your interests in your area or online



[aka.ms/SuperUsers](https://aka.ms/SuperUsers)

Spread your Fabric knowledge, insights, and best practices with others



[aka.ms/MVP](https://aka.ms/MVP)

Technology experts that share their knowledge and passion with the community

Community Conference



Restrooms

Breaks and lunch

Questions

Workshop content

What is your learning style?



# Agenda

01

Introduction to Fabric

05

Data Transformation  
and T-SQL

09

Security

02

Introduction to Data  
Warehouse

06

Monitoring Warehouse  
Activity

03

Creating Tables

07

Statistics and V-Order

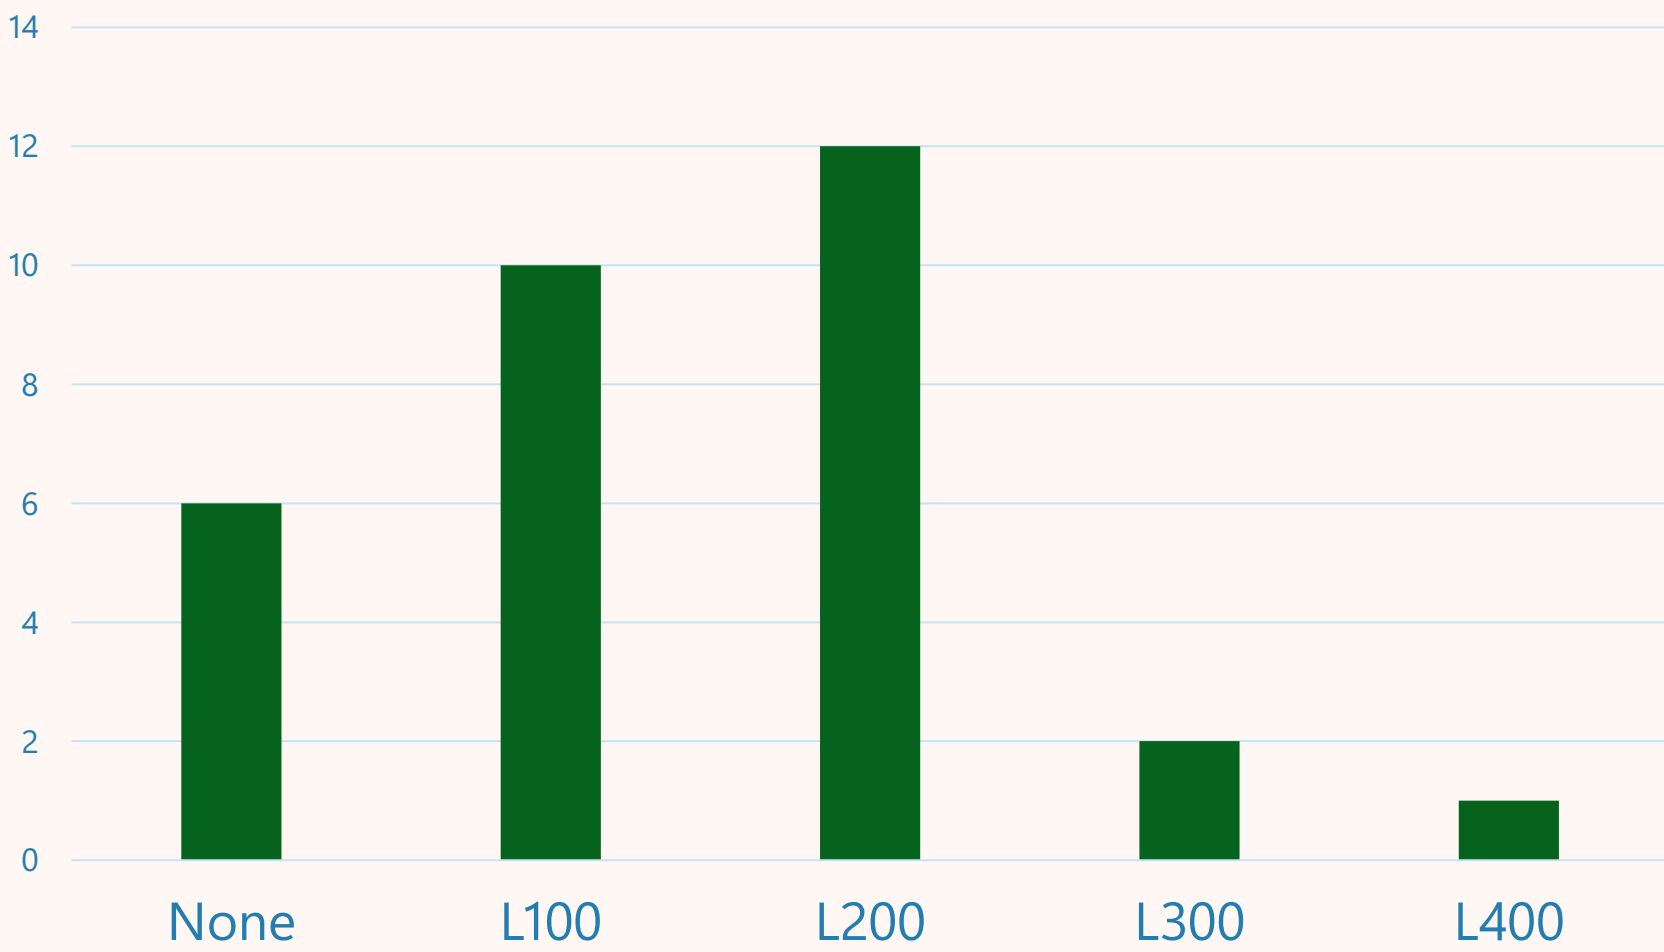
04

Data Ingestion

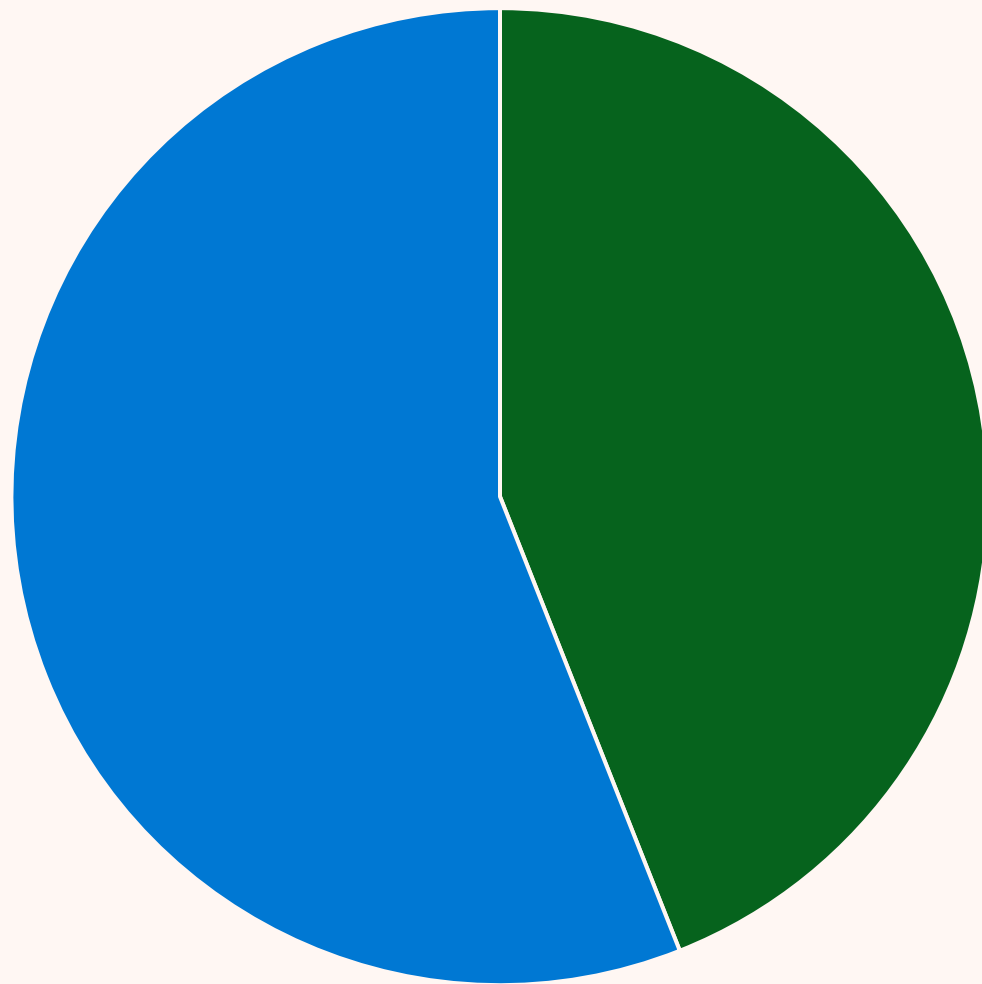
08

Workload  
Management

# Poll: What level of experience do you have with Fabric Data Warehouse?

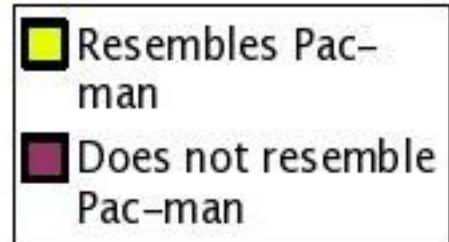
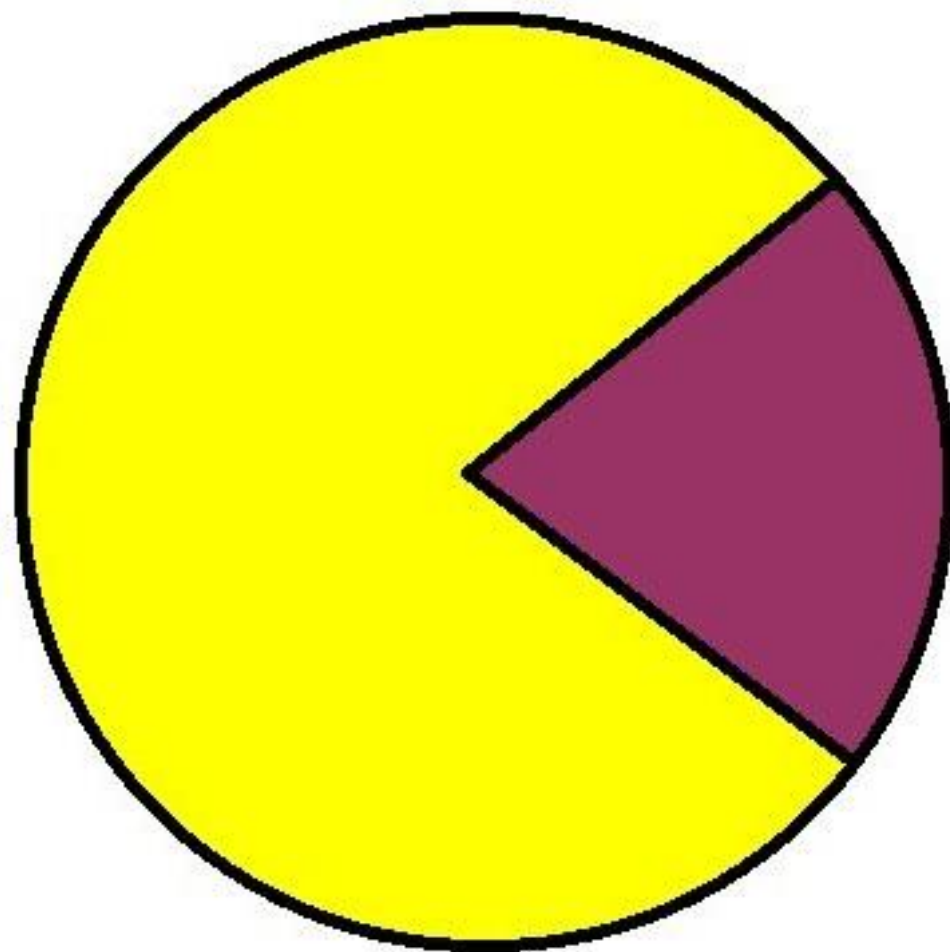


# Poll: Should pineapple be on pizza?



■ No    ■ Are you insane?! Absolutely not!

## Percentage of Chart Which Resembles Pac-man



Community Conference



# Introduction to Fabric

Module 1

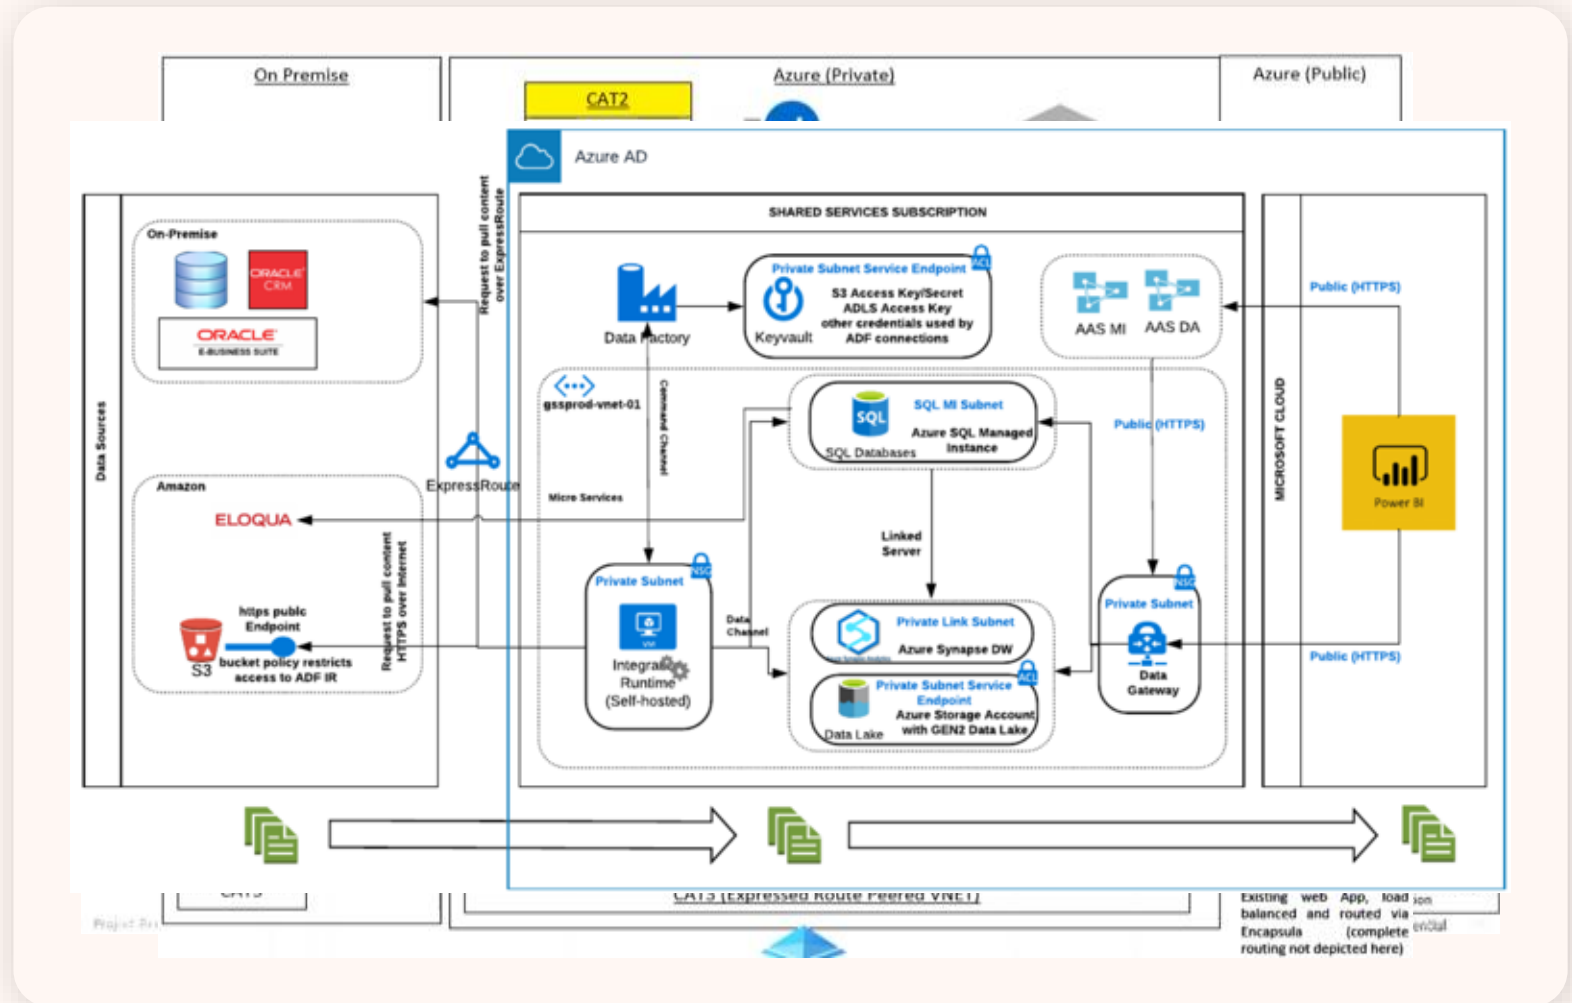
# Analytics is complex and fragmented

Every project has many subsystems

Every subsystem need a different class of product

Products often comes from multiple vendors

Integration is complex, fragile and expensive





# Microsoft Fabric



Data  
Factory



Real-Time  
Intelligence



Databases



Data  
Warehouse



Data  
Engineering



Power BI



Industry  
Solutions



Copilot in Fabric



OneLake



Microsoft Purview

# The data platform for the era of AI

## Complete Analytics Platform

Unified product, experience, and architecture

---

Delivered as SaaS

## Lake Centric and Open

Common SaaS data lake shared by all compute engines

---

Deep commitment for open formats and APIs

## Empower Every Business User

Deliver data directly to the users in their favorite Office applications

## AI Powered

Designed from the ground up for AI

# Microsoft Fabric



Data  
Factory



Real-Time  
Intelligence



Databases



Data  
Warehouse



Data  
Engineering



Power BI



Industry  
Solutions



Copilot in Fabric



OneLake

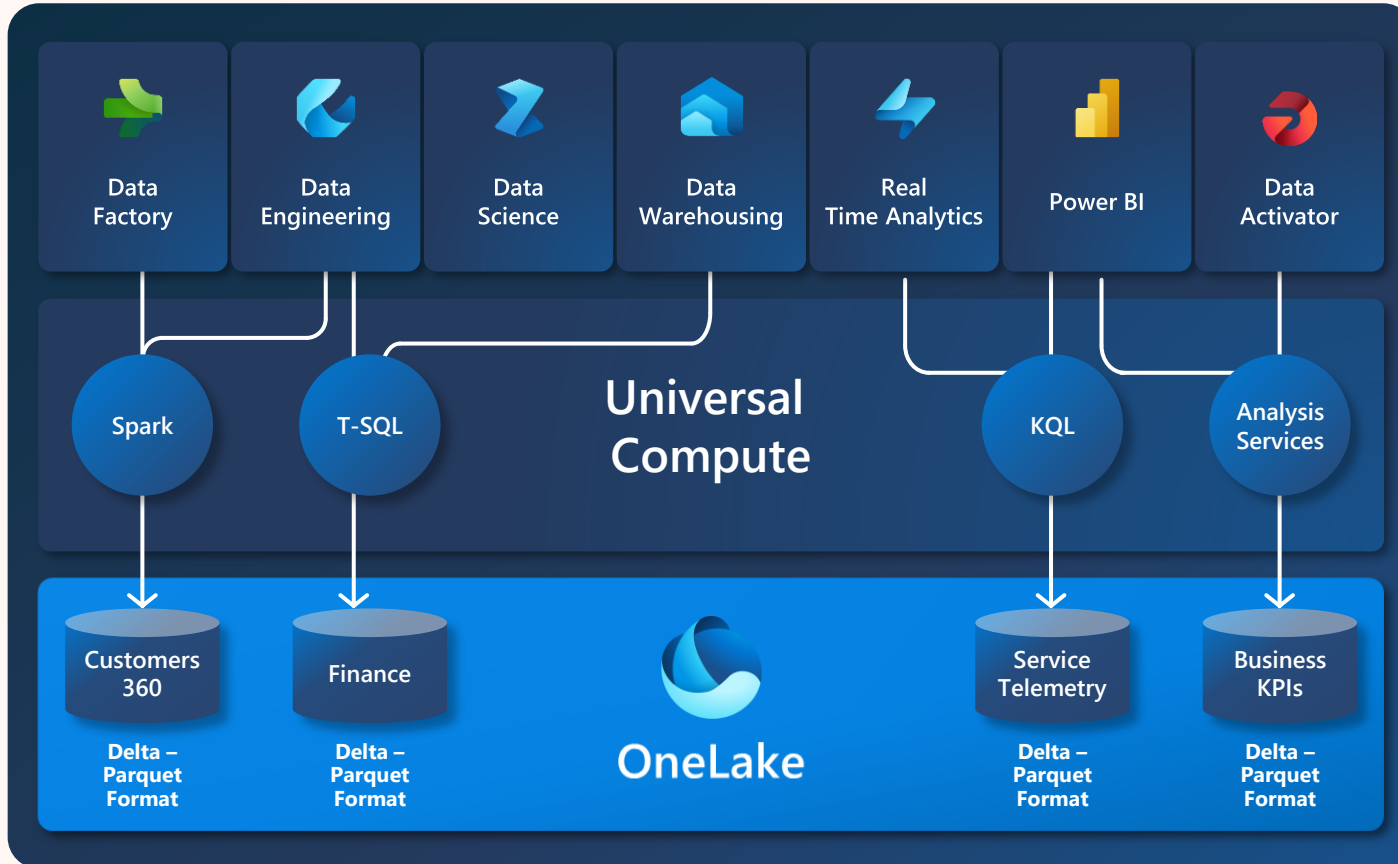


Microsoft Purview

## Single...

Onboarding and trials  
Sign-on  
Navigation model  
UX model  
Workspace organization  
Collaboration experience  
Data Lake  
Storage format  
Data copy for all engines  
Security model  
CI/CD  
Monitoring hub  
Data Hub  
Governance & compliance

# Microsoft Fabric



All the compute engines store their data automatically in OneLake

The data is stored in a single common format

**Delta - Parquet**, an open standards format, is the storage format for all tabular data in Analytics vNext

Once data is stored in the lake, it is directly accessible by all the engines without needing any import/export

All the compute engines have been fully optimized to work with Delta Parquet as their native format


Shared universal security model is enforced across all the engines

# Fabric topology and organization

# Fabric Tenant

## Capacity 1 – North Europe

### Workspace 1



Data  
Factory



Synapse Data  
Engineering

## Capacity 2 – East US

### Workspace 3




Synapse Data  
Science

## Capacity 3 – East US


### Workspace 2



Synapse Data  
Warehousing



Power BI



Data  
Activator


### Workspace 4



Synapse Data  
Engineering



Synapse Data  
Warehousing



Power BI

OneLake



# Comparing P and F SKUs

SKU	Capacity Units (CU)	Power BI SKU	Power BI v-cores
F2	2	-	0.25
F4	4	-	0.5
F8	8	EM/A1	1
F16	16	EM2/A2	2
F32	32	EM3/A3	4
<b>F64</b>	<b>64</b>	<b>P1/A4</b>	<b>8</b>
F128	128	P2/A5	16
F256	256	P3/A6	32
F512	512	P4/A7	64
F1024	1024	P5/A8	128
F2048	2048	-	256

# Getting started with Fabric

# Getting started in your environment

1. Enable Fabric at the tenant level.
  - Override the tenant level setting at the capacity
  - Enable for only specific security groups
2. Get capacity
  - Start a 60-day free trial
  - Use existing P SKU
  - Provision an F SKU
3. Create a workspace and assign a capacity.
4. Build something cool!

Users can create Fabric items

*Enabled for the entire organization*

Users can use production-ready features to create Fabric items. Turning off this setting doesn't impact users' ability to create Power BI items. This setting can be managed at both the tenant and the capacity levels. [Learn More](#)



Enabled

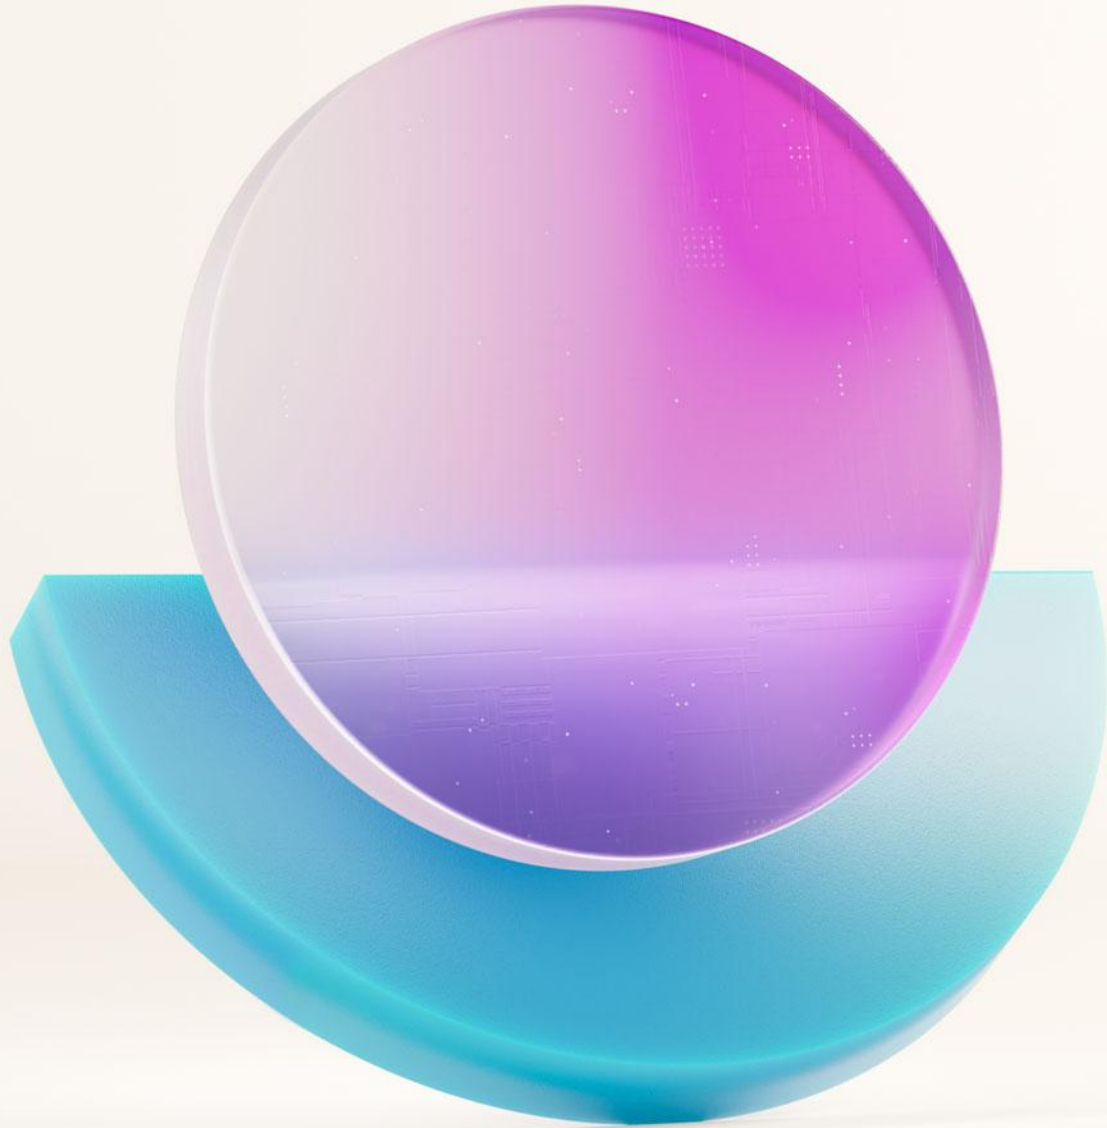
Users can try Microsoft Fabric paid features

*Enabled for the entire organization*

When users sign up for a Microsoft Fabric **trial**, they can try Fabric paid features for free for 60 days from the day they signed up. [Learn More](#)



Enabled



# Introduction to Data Warehouse

## Module 2

# Microsoft Fabric



Data  
Factory



Real-Time  
Intelligence



Databases



Data  
Warehouse



Data  
Engineering



Power BI



Industry  
Solutions



Copilot in Fabric



OneLake



Microsoft Purview

# What is a data warehouse?

A data warehouse is a central repository of information that can be analyzed to make more informed decisions. Data flows into a data warehouse from transactional systems, relational databases, and other sources, typically on a regular cadence. Business analysts, data engineers, data scientists, and decision makers access the data through business intelligence tools, SQL clients, and other analytics applications.



# Data Warehouse key capabilities

Truly serverless data warehouse

Data warehouse redesigned from the ground up

Optimized for star schema designs

Familiar T-SQL code

Easy migration path

Use it with a lakehouse or standalone

Auto stats

Auto file maintenance

No indexes required

SQL Best Practices still apply

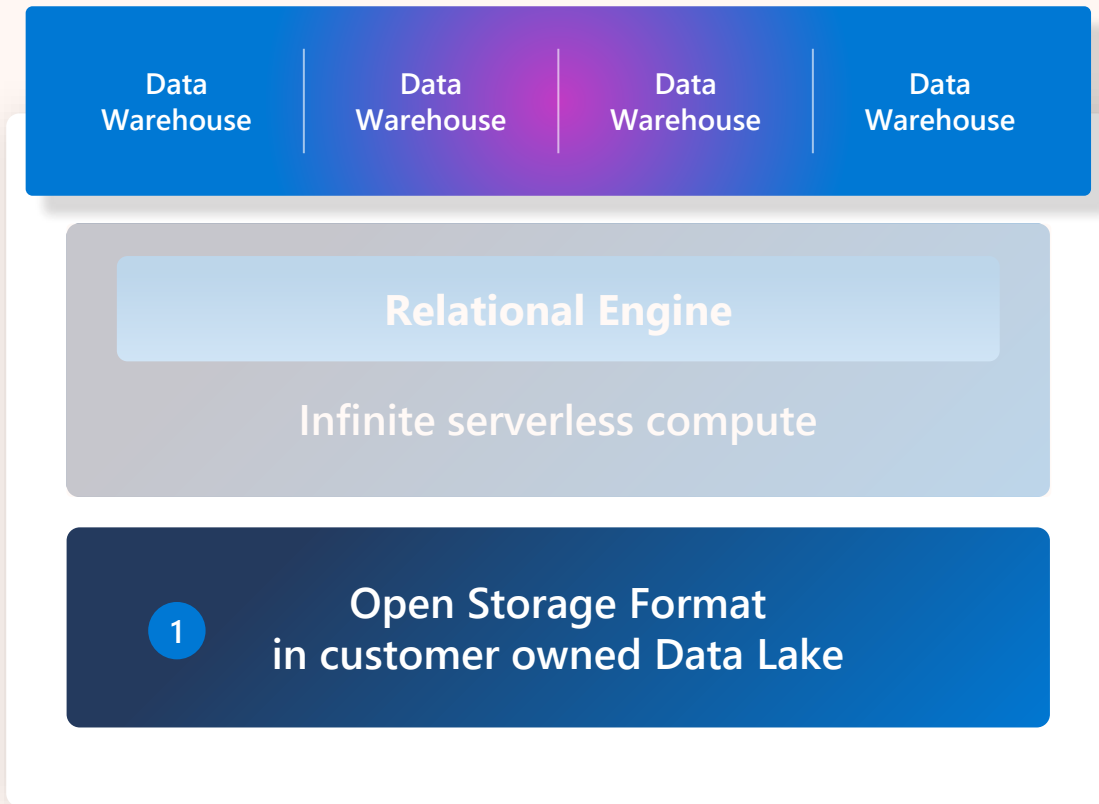
Rapid innovation (release cycle, new features, closing gaps)

# Fabric Data Warehouse

Infinitely scalable and open



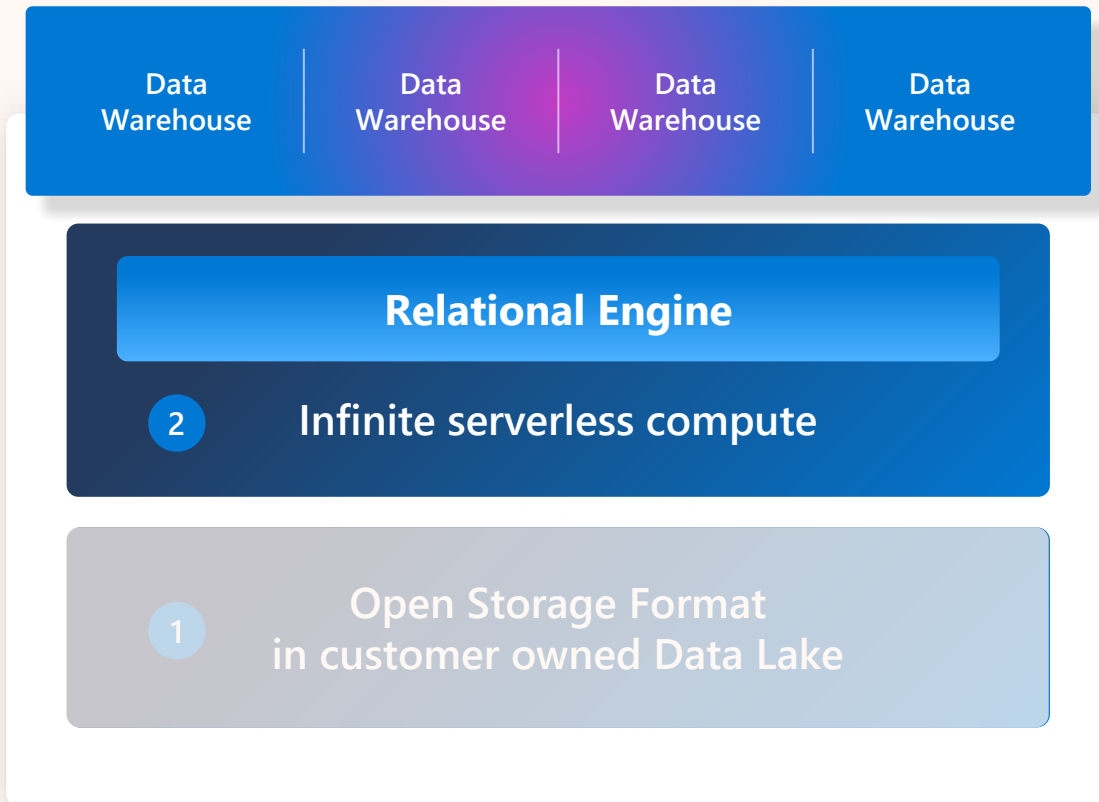
## Fabric Data Warehouse



- 1 Open standard format in an open data lake replaces proprietary formats as the native storage
  - First transactional data warehouse natively embracing an open standard format
  - Data is stored in Delta – Parquet with no vendor lock-in
  - Is auto-integrated and auto-optimized with minimal knobs
  - Extends full SQL ecosystem benefits

# Fabric Data Warehouse

Infinitely scalable and open



## 2 Dedicated clusters are replaced by serverless compute infrastructure

- Physical compute resources assigned within milliseconds to jobs
- Infinite scaling with dynamic resource allocation tailored to data volume and query complexity
- Instant scaling up/down with no physical provisioning involved
- Resource pooling providing significant efficiencies and pricing

# Data Modeling



**Star schema** is the recommended design approach to take when creating a Fabric warehouse.



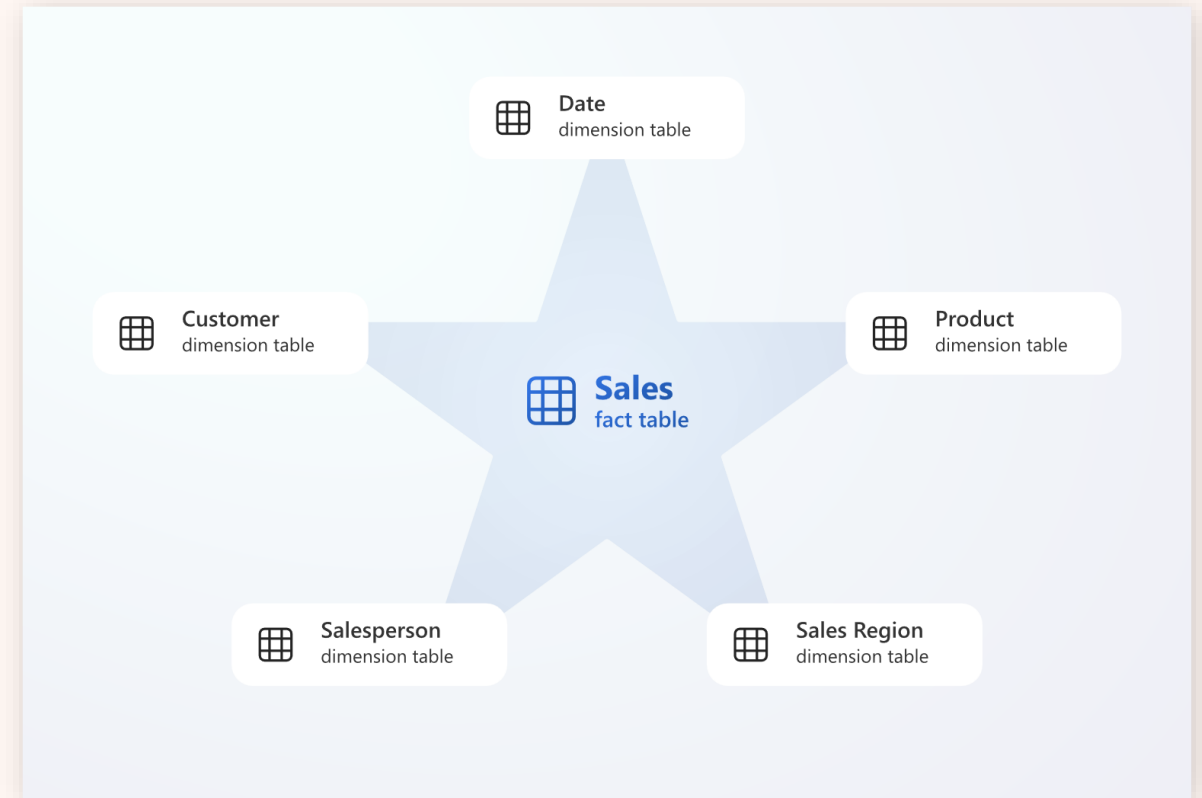
An **OLTP** workload can run on Fabric warehouse but is not recommended for optimal performance.



The warehouse engine is optimized for **batch** writes and reads.



Micro-batches can work well, but **trickle feeds** should be avoided where possible.



# Dimensional Modeling

## Facts tell us...

What happened

When did it happen

Contain measures (counts, quantities, amounts, etc.)

Don't typically change

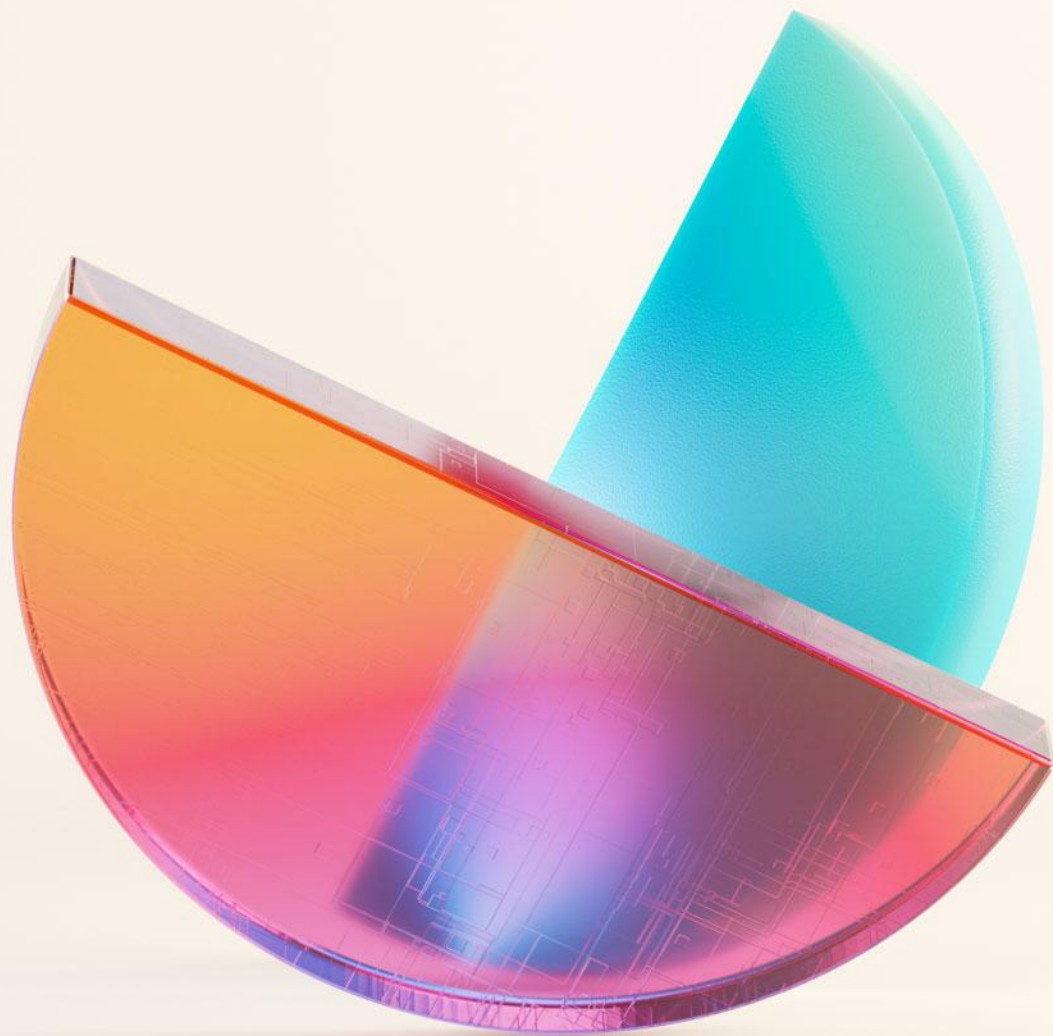
## Dimensions describe...

What happened in a "fact"

Ways to slice and dice an event

Attributes of the "member"

How attributes change over time



# Creating Tables

## Module 3



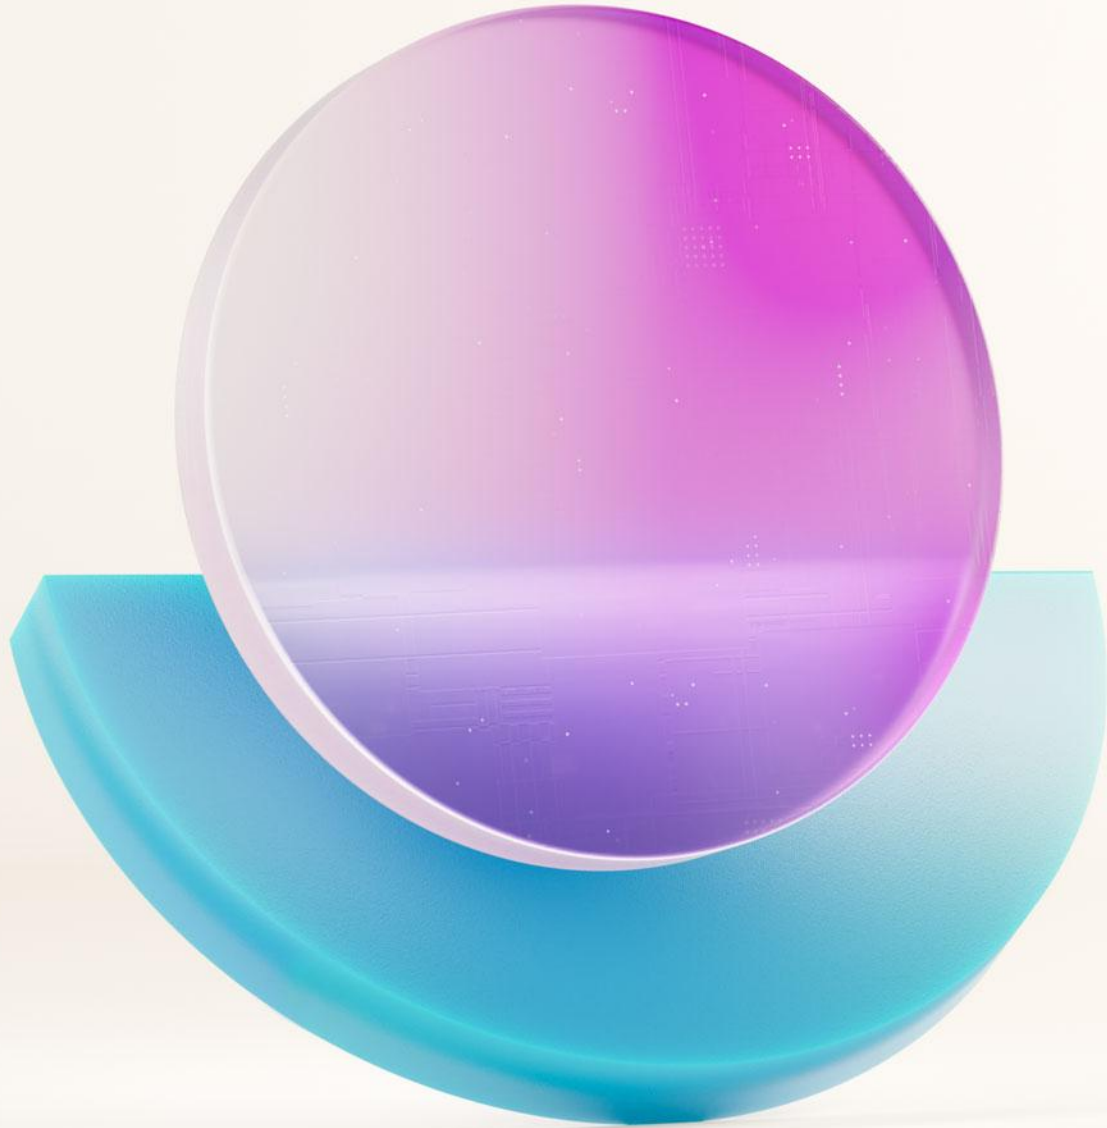
# Creating Data Warehouse tables

Basic syntax:

```
CREATE TABLE { database_name.schema_name.table_name | schema_name.table_name | table_name }  
    (  
        { column_name <data_type> [ <column_options> ] } [ ,...n ]  
    )
```

Example:

```
CREATE TABLE Bands (  
    BandID INT NOT NULL,  
    Name CHAR(45),  
    FoundationDate DATE  
)
```



Demo

# Data type support

Exact numeric	Approximate numeric	Date and Time	Character strings	Binary strings	Other
bit bigint int smallint decimal(p,s) numeric(p,s)	float(n) real	date datetime2(n)* time(n)*	char(n) varchar(n) char(max)** varchar(max)**	varbinary(n)	uniqueidentifier

\* datetime2(6) and time(6) are the limits

\*\* char/varchar(max) currently in preview

# Optimizing data types for better performance

Best practices to maximize performance:

## **Use the appropriate data type**

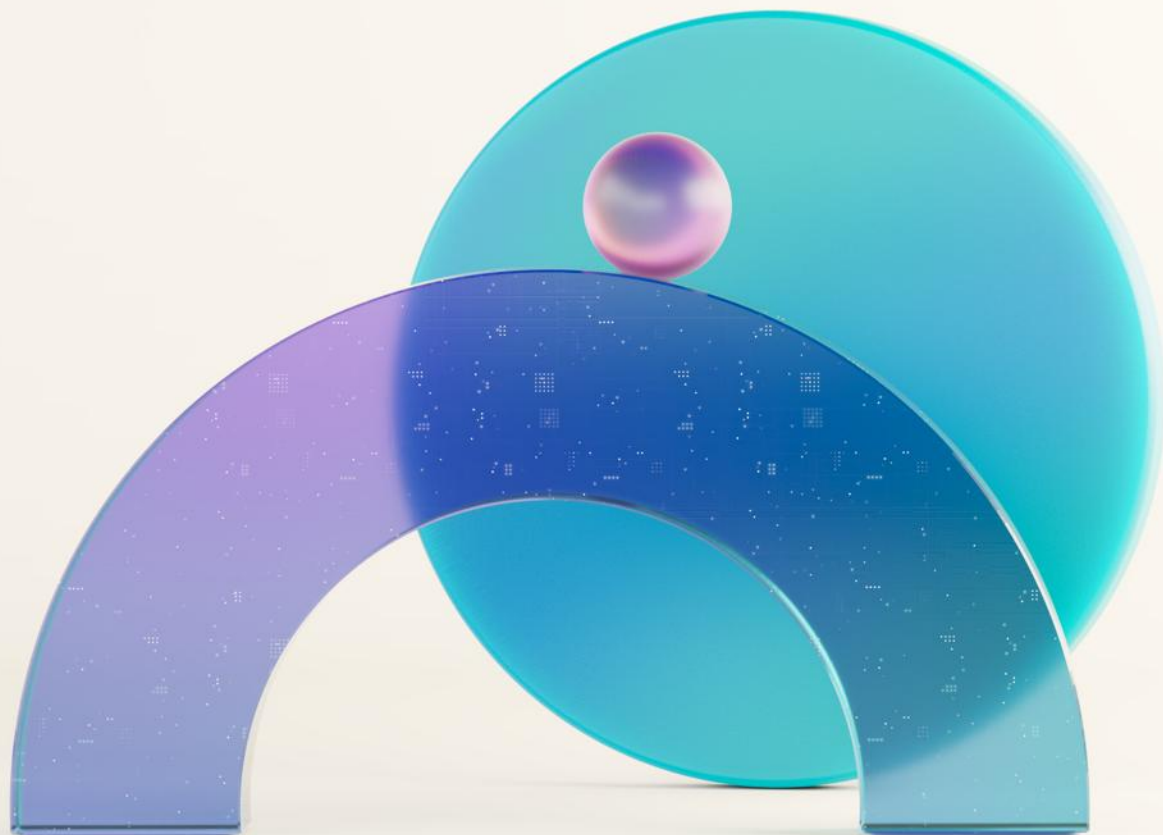
- If a column contains only integer values, use INT data types instead of DECIMAL/NUMERIC

## **Precision matters**

- Reduce DECIMAL/NUMERIC precision to the smallest one that can accommodate your data
- Don't default to CHAR/VARCHAR(8000). Adjust the precision to your requirements (e.g.: VARCHAR(50))

## **Always use the smallest data type that supports your data**

- Prefer INT over BIGINT when possible, or SMALLINT over INT
- Prefer DATE or TIME over DATETIME2 if you only need the date or time alone

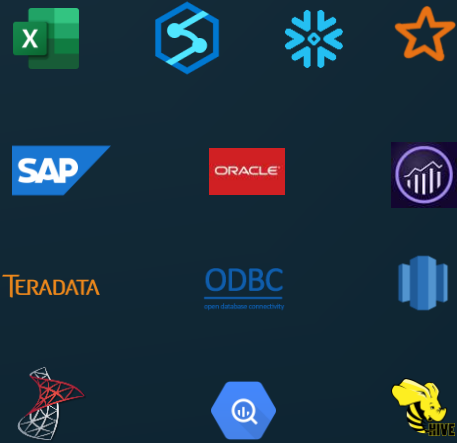


# Data Ingestion

## Module 4

# Data orchestration and data prep with Pipelines

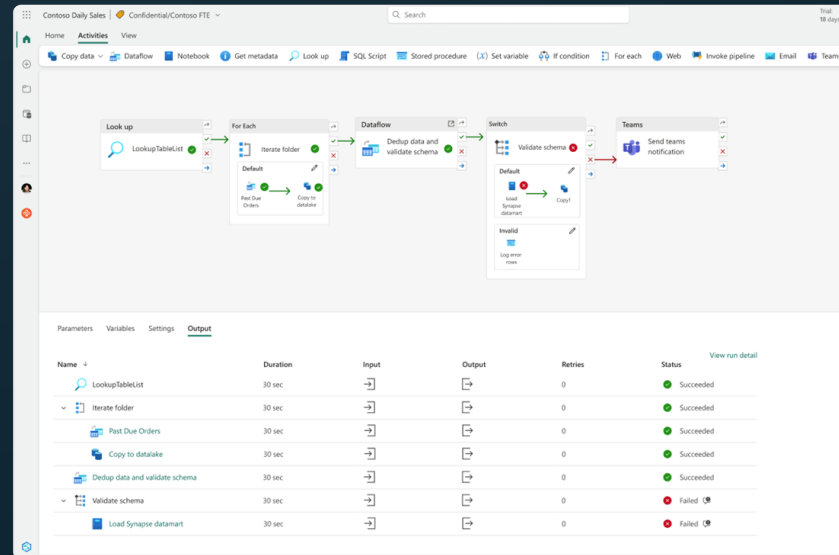
## The worlds data



Through 170+ connectors



## ETL, CDC, or ELT data prep needs



Rich library of activities and pro code data prep experiences

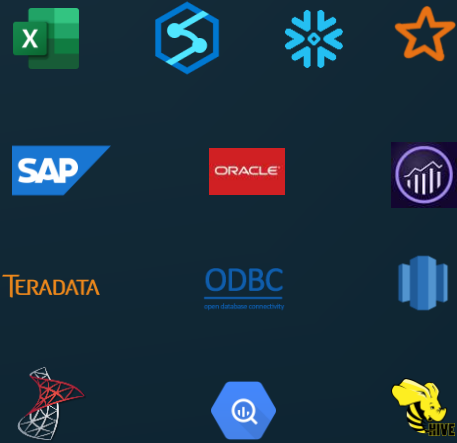


Warehouse

Ingest and transform data at scale

# No code data prep with Dataflows gen 2

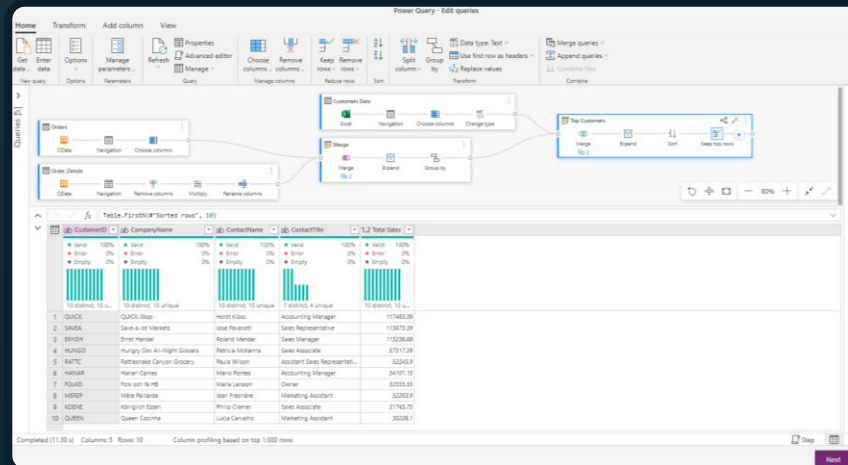
## The worlds data



Through 170+ connectors



## Familiar Power Query experiences



Through 400+ no code transformations



Warehouse

Ingest and transform data at scale

# ...and of course, you can use T-SQL too!

## SQL COPY INTO, transactions, and more

Support for:	COPY INTO
	CTAS
	INSERT ... SELECT
	Common Table Expressions
	Stored Procedures
Transactions and locking:	Supports snapshot isolation
	INSERTS and SELECTS only require SCHEMA stability locks
	UPDATES and DELETES do not block tables

## Simple, Familiar SQL Syntax

```
COPY INTO table_name.[(Column_list)]
FROM '<external_location>' [,...n]
WITH
(
  [FILE_TYPE = { 'CSV' | 'PARQUET' }]
  [, CREDENTIAL = (AZURE CREDENTIAL)]
  [, COMPRESSION = { 'Gzip' | 'Snappy' }]
  [, FIELDQUOTE = 'string_delimiter']
  [, FIELDTERMINATOR = 'field_terminator']
  [, ROWTERMINATOR = 'row_terminator']
  [, FIRSTROW = first_row]
  [, ENCODING = { 'UTF8' | 'UTF16' }]
);

CREATE TABLE [dbo].[FactInternetSales_new]
AS
SELECT * FROM [dbo].[FactInternetSales];

INSERT INTO dbo.EmployeeSales
SELECT 'SELECT', sp.BusinessEntityID, c.LastName, sp.SalesYTD
FROM warehouseA.Sales.SalesPerson AS sp
INNER JOIN warehouseB.Person.Person AS c
ON sp.BusinessEntityID = c.BusinessEntityID
WHERE sp.BusinessEntityID LIKE '2%'
ORDER BY sp.BusinessEntityID, c.LastName;

WITH sales_data AS (
  SELECT customer_id, SUM(total_sales) AS total_sales
  FROM sales
  GROUP BY customer_id
)
SELECT customer_id, total_sales
FROM sales_data
WHERE total_sales > 1000;
```



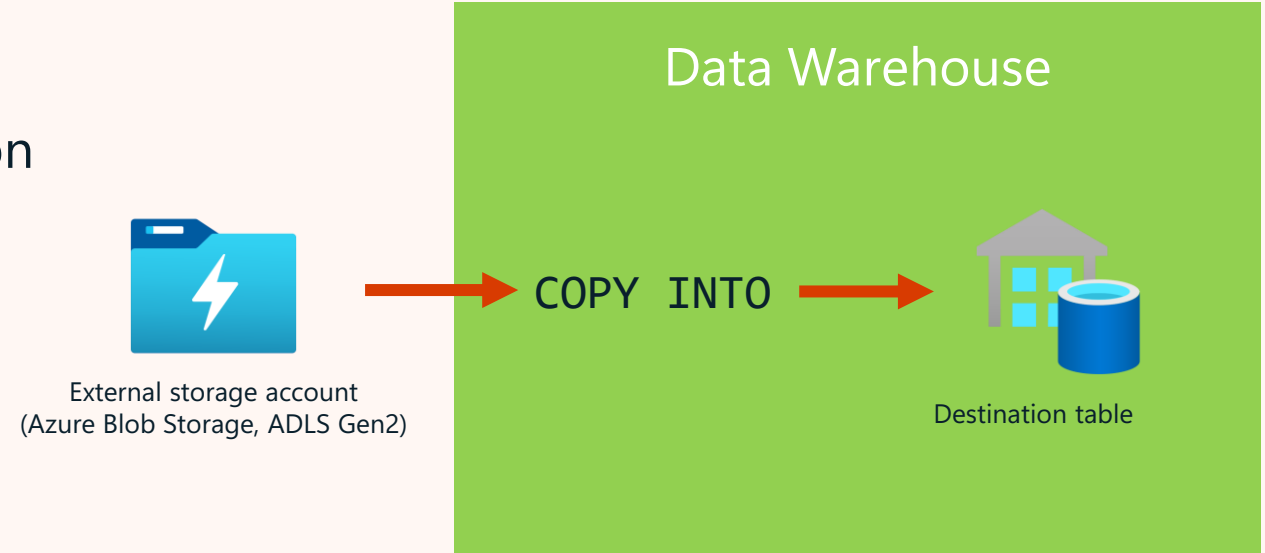
# COPY INTO

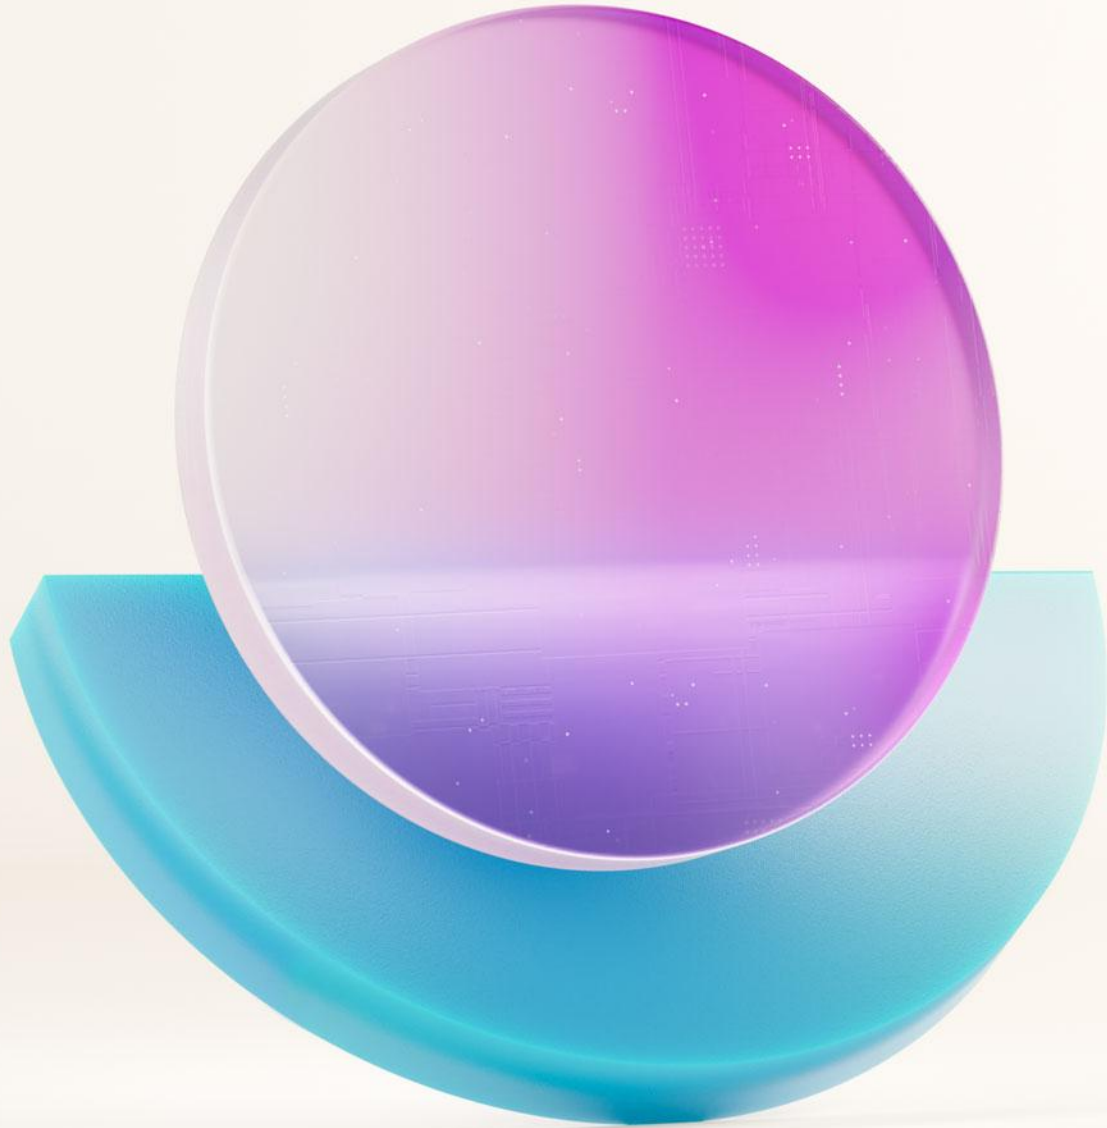
Enables flexible, high-throughput data ingestion from external Azure storage accounts

- Supports PARQUET and CSV formats
- No additional database objects required

Example:

```
COPY INTO SalesOrderItems
FROM 'https://myaccount.blob.core.windows.net/mycontainer/*.csv'
WITH (
    FILE_TYPE = 'CSV',
    FIRSTROW = 2
)
```





Demo

# Ingesting data from tables on other warehouses

Cross-warehouse queries empower seamless integration of data from multiple warehouses, enabling comprehensive analysis, actionable insights, and efficient workflows for improved productivity.

- **Unified Data Analysis:** Combine datasets from different warehouses to gain a holistic view of business operations.
- **Enhanced Insights:** Identify trends and correlations across isolated datasets to drive informed decisions.
- **Boosted Productivity:** Streamline workflows by accessing and analyzing data from multiple sources in a single query.

# Ingesting data from tables on other warehouses (cont'd)

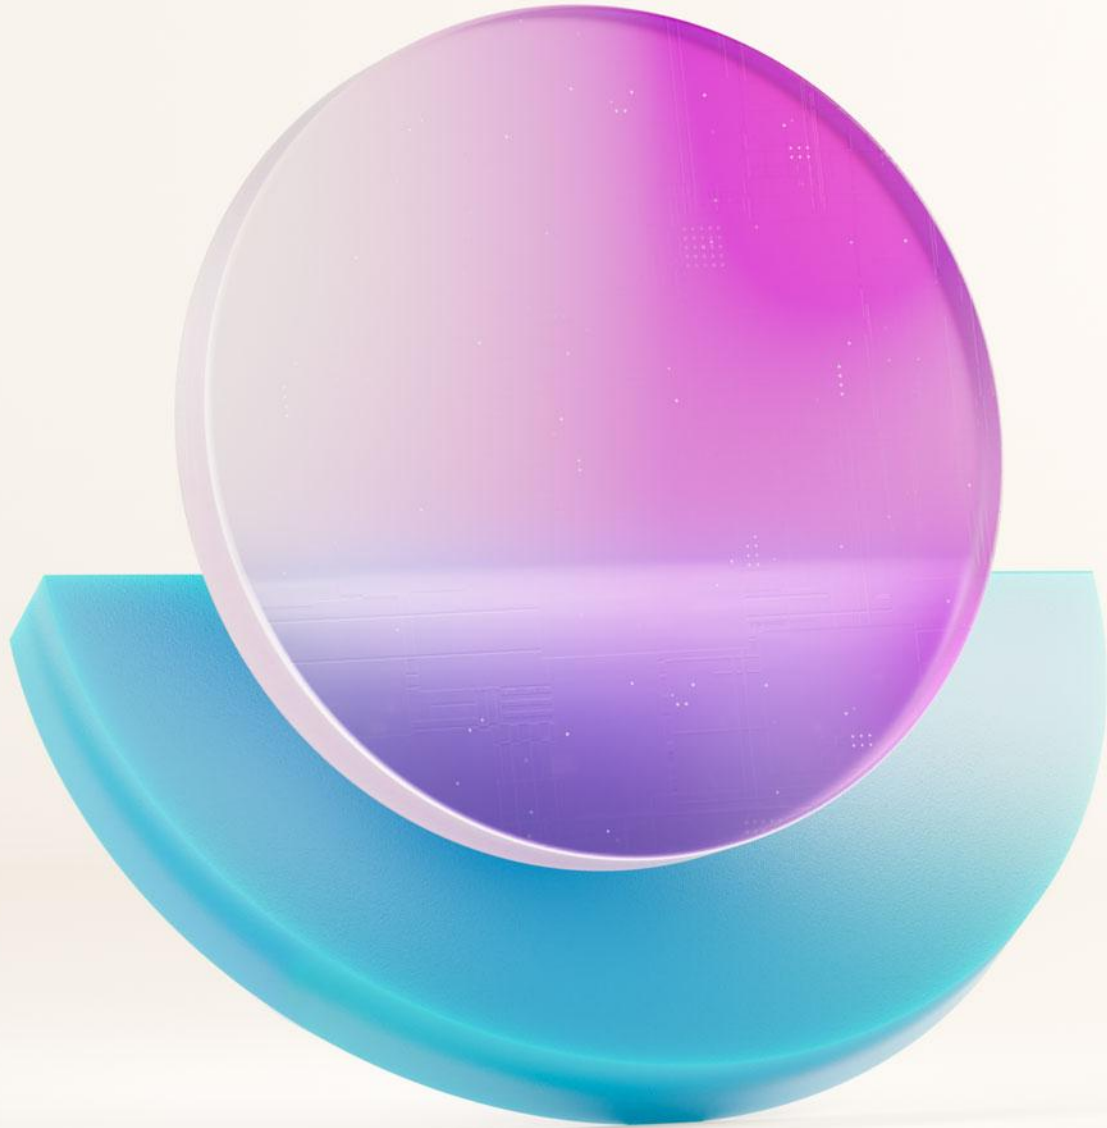
Cross-warehouse queries are allowed using three-part naming:

```
SELECT <column list>  
FROM <WarehouseName>.<SchemaName>.<TableName>
```

Example using CREATE TABLE AS SELECT (CTAS):

```
CREATE TABLE WarehouseC.Analytics.CustomerSalesSummary AS  
SELECT  
    WHa.CustomerName,  
    WHb.OrderDate,  
    WHb.TotalOrderValue  
FROM WarehouseA.Customers.CustomerData AS WHa  
JOIN WarehouseB.Orders.SalesData AS WHb  
ON WHa.CustomerID = WHb.CustomerID  
WHERE WHb.TotalOrderValue > 100;
```

Other T-SQL patterns are also supported (e.g.: SELECT INTO, INSERT INTO... SELECT)



Demo

# Mirroring

Mirroring in Fabric is a low-cost and low-latency solution to bring data from various systems together into a single analytics platform. You can continuously replicate your existing data estate directly into Fabric's OneLake from a variety of Azure databases and external data sources.

Enabling Mirroring is simple and intuitive, without having the need to create complex ETL pipelines, allocate other compute resources, and manage data movement.

It is a fully managed service, so you don't have to worry about hosting, maintaining, or managing replication of the mirrored connection.

Platform	Near real-time replication	Type of mirroring
<a href="#">Microsoft Fabric mirrored databases from Azure Cosmos DB (preview)</a>	Yes	Database mirroring
<a href="#">Microsoft Fabric mirrored databases from Azure Databricks (preview)</a>	Yes	Metadata mirroring
<a href="#">Microsoft Fabric mirrored databases from Azure SQL Database</a>	Yes	Database mirroring
<a href="#">Microsoft Fabric mirrored databases from Azure SQL Managed Instance (preview)</a>	Yes	Database mirroring
<a href="#">Microsoft Fabric mirrored databases from Snowflake</a>	Yes	Database mirroring
<a href="#">Open mirrored databases (preview)</a>	Yes	Open mirroring
<a href="#">Microsoft Fabric mirrored databases from Fabric SQL database (preview)</a>	Yes	Database mirroring

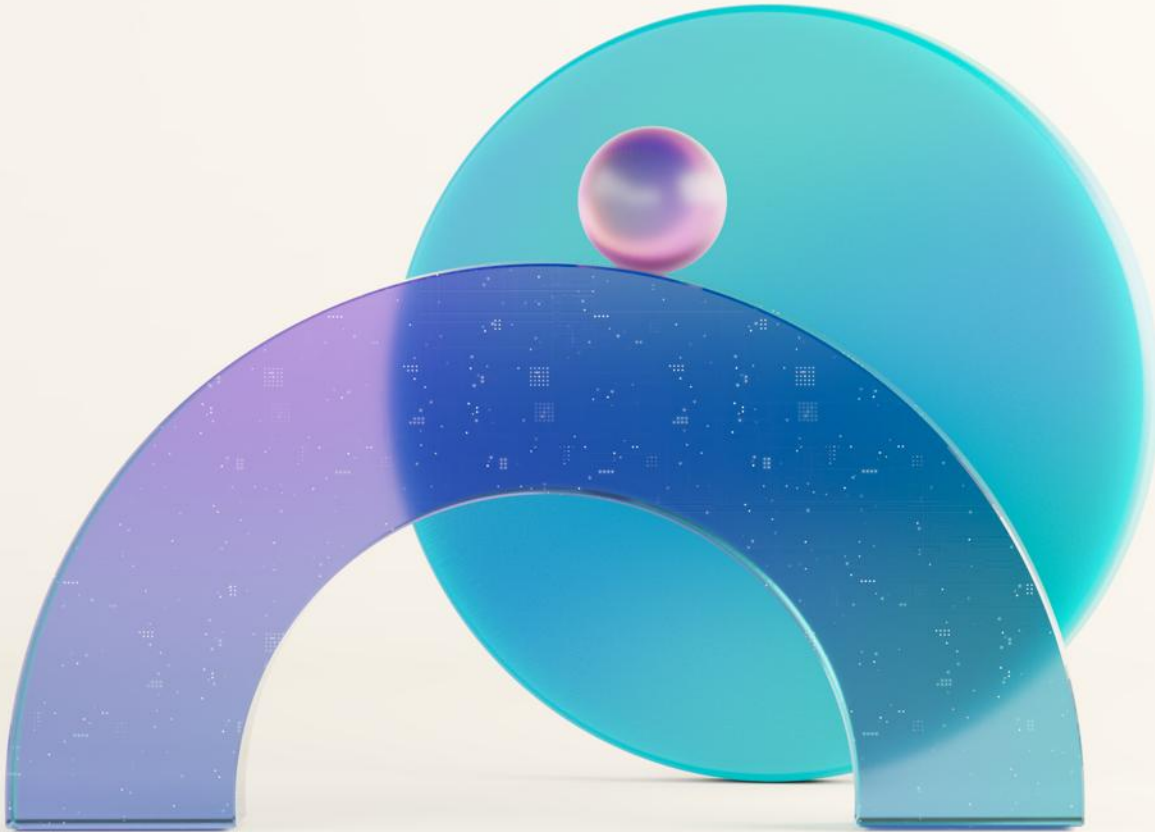
Availability of Mirroring by platform



10-minute break

# Lunch

## We will start at 1:00 PM







# Data Transformation and T-SQL

Module 5

# How do you perform data transformation?

## At Load Time

Extract Transform Load

Dataflows Gen2

Spark

SQL Server Integration Services (SSIS)

## After Loading

Extract Load Transform

T-SQL Scripts (at query time)

Stored Procedures

Functions

# Programmability

## Stored procedures

Encapsulate sets of logic written in T-SQL.

Can be called through T-SQL or through Data Factory pipeline activities.

Run any T-SQL statement in the supported surface area.

## Functions

Support for table-valued functions

Scalar functions coming soon

Repeatable logic that can be used independently or referenced in other T-SQL statements

# Common Data Warehouse Operations



## Surrogate Keys

Synthetic key only used in the data warehouse

Identity support on the roadmap

```
DECLARE @MaxKey  
SELECT MAX(Key) FROM Table  
  
...@MaxKey + RowNumber...
```



## Type 1 and 2 Attributes

MERGE coming soon!

### Update

Type 1 attributes  
End date type 2 records

### Insert

Insert new records  
Insert new type 2 versions



## Late Arriving Members

Create "unknown" members in dimension

Best practice: use -1 or 0 for the unknown member

# Time Travel

Historical analysis of data up to 30 days in the past.

Granularity of the time travel is at the statement level. Cannot vary times for each table within a query.

Audit changes for compliance purposes.

Only returns the latest version of the table schema.

Dropping a table removes the history.

## T-SQL Syntax

**SELECT**

[Columns]

**FROM** [schema].[Table or View]

[**JOIN** Table or View]

**OPTION** (**FOR TIMESTAMP AS OF** 'YYYY-MM-DDTHH:MM:SS.ss');

# Cloning tables

“Copy” data with zero duplication of data.

Create historical reports reflecting data as it was at a specific point in time.

Clones become a new branch of the table allowing for full, independent DDL and DML without impact to the original table.

Security is inherited by the clone.

## T-SQL Syntax

```
CREATE TABLE { database_name.schema_name.table_name | schema_name.table_name | table_name }  
AS CLONE OF { database_name.schema_name.table_name | schema_name.table_name | table_name }  
[AT {point_in_time}]
```

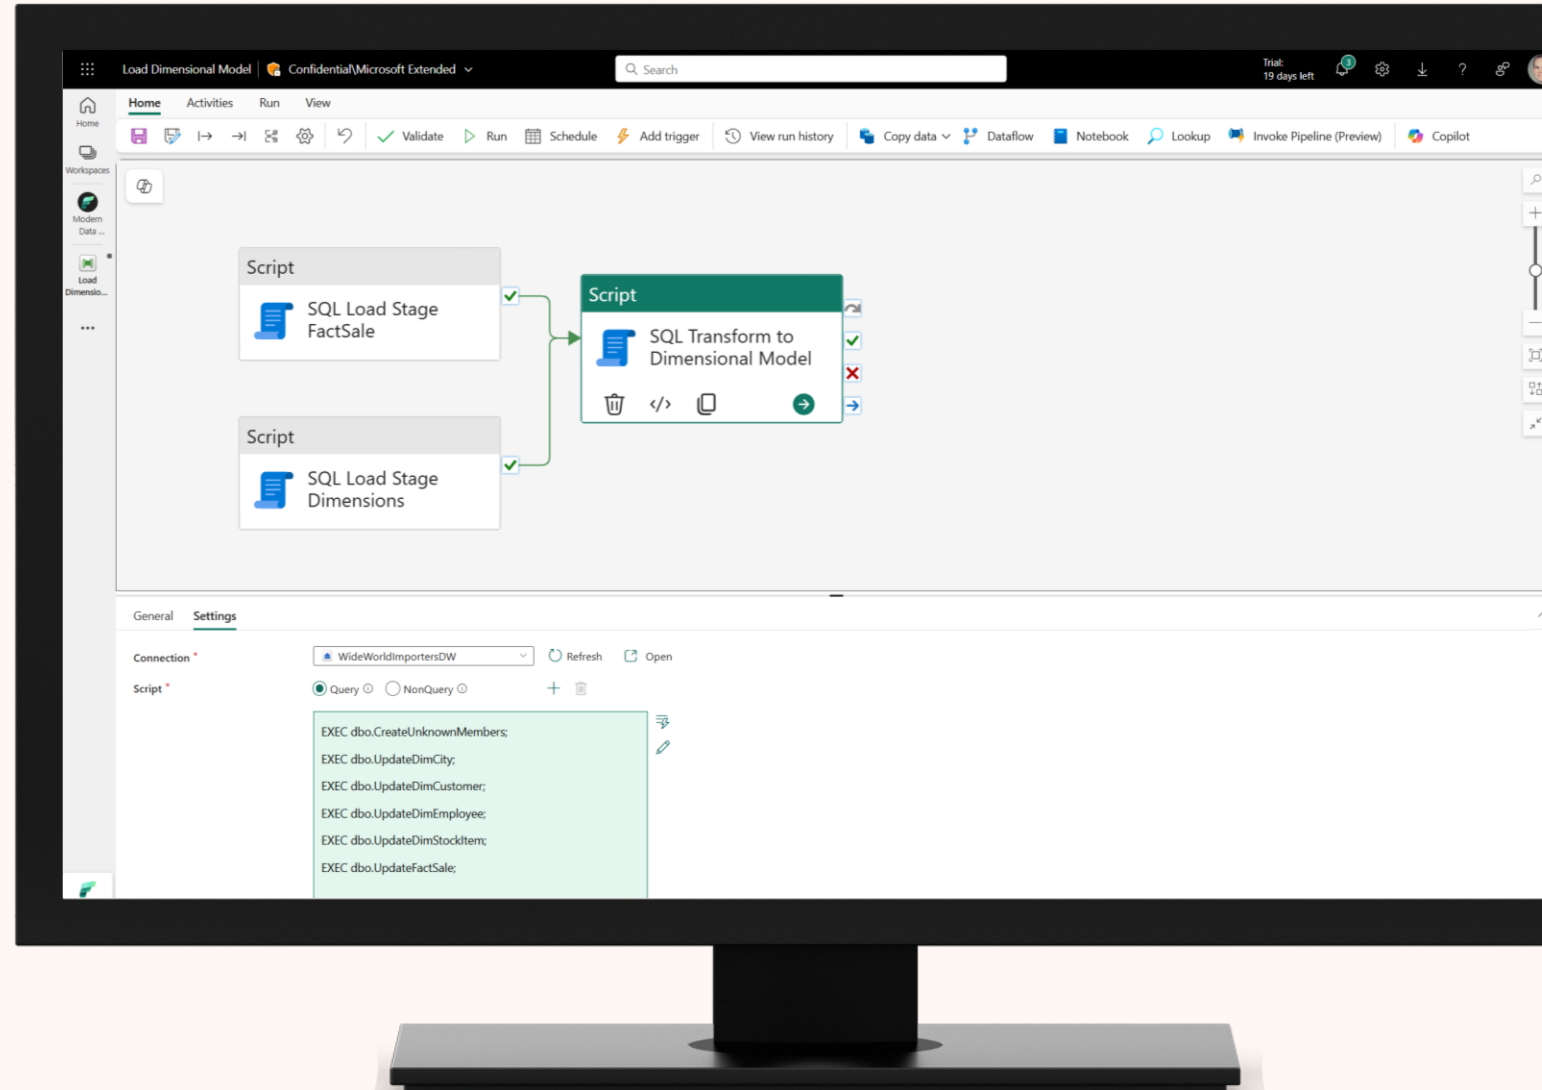
# Orchestrating ETL Operations with Pipelines

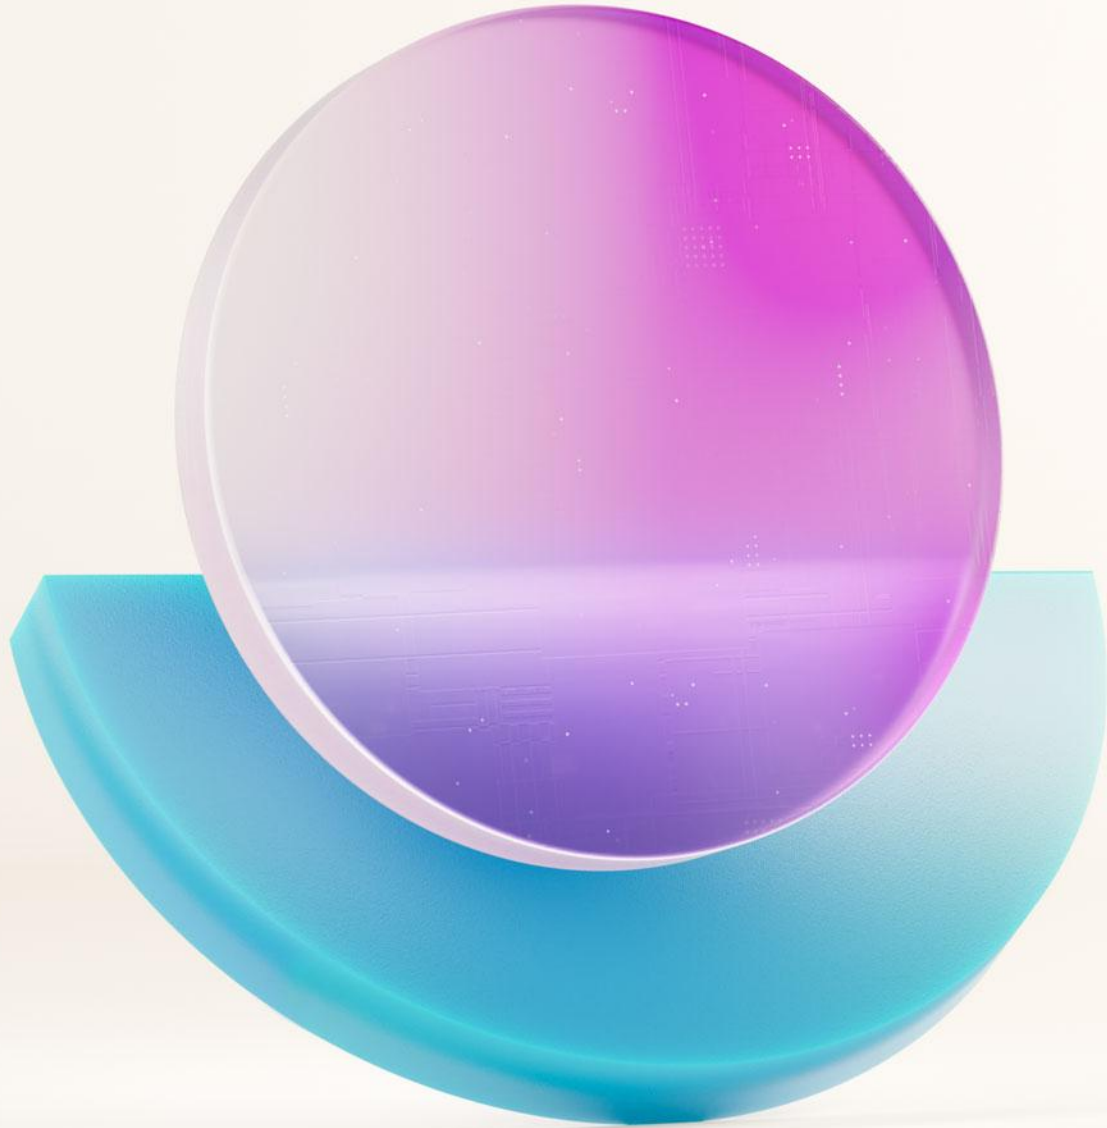
Data Factory pipelines move data or dispatch operations to other compute engines.

Ensure proper order of operations.

Facilitate scheduling ETL operations.

Can be simple “run this code” or complex metadata driven frameworks.





# Monitoring Warehouse Activity

Module 6



# Dynamic Management Views

## **sys.dm\_exec\_connections**

Information about each connection established to the engine.

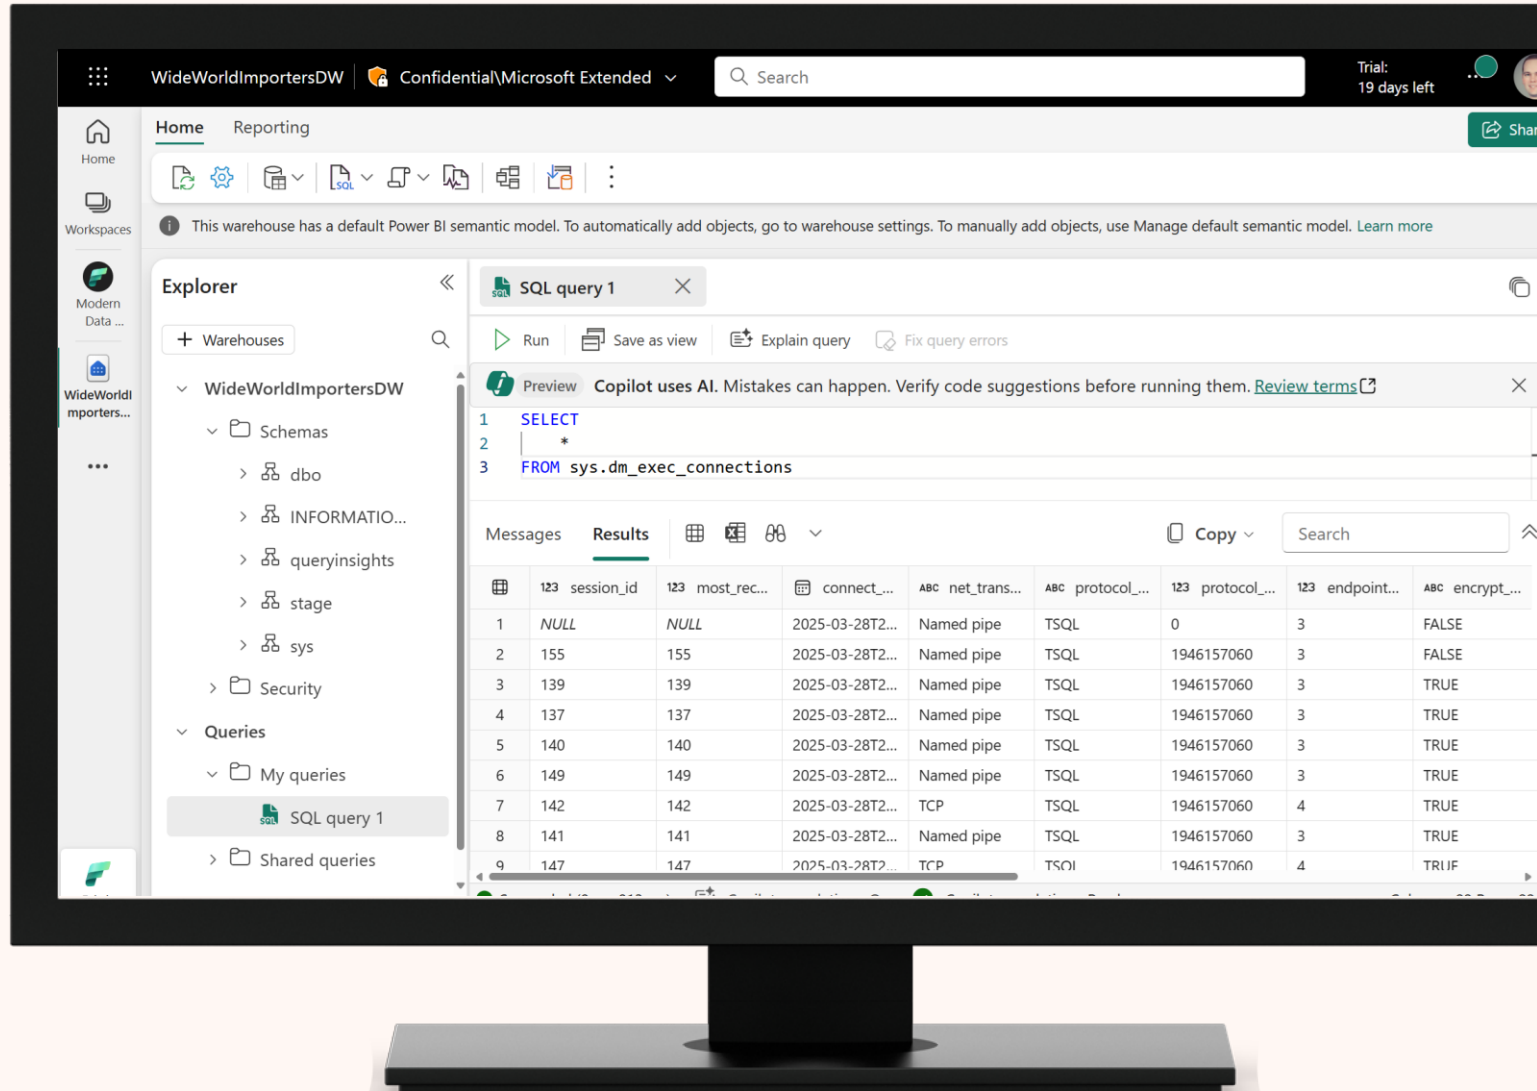
## **sys.dm\_exec\_sessions**

Information about each session.

## **sys.dm\_exec\_requests**

Information about each request made by users and generated by the system.

Member, Contributor, or View permissions.



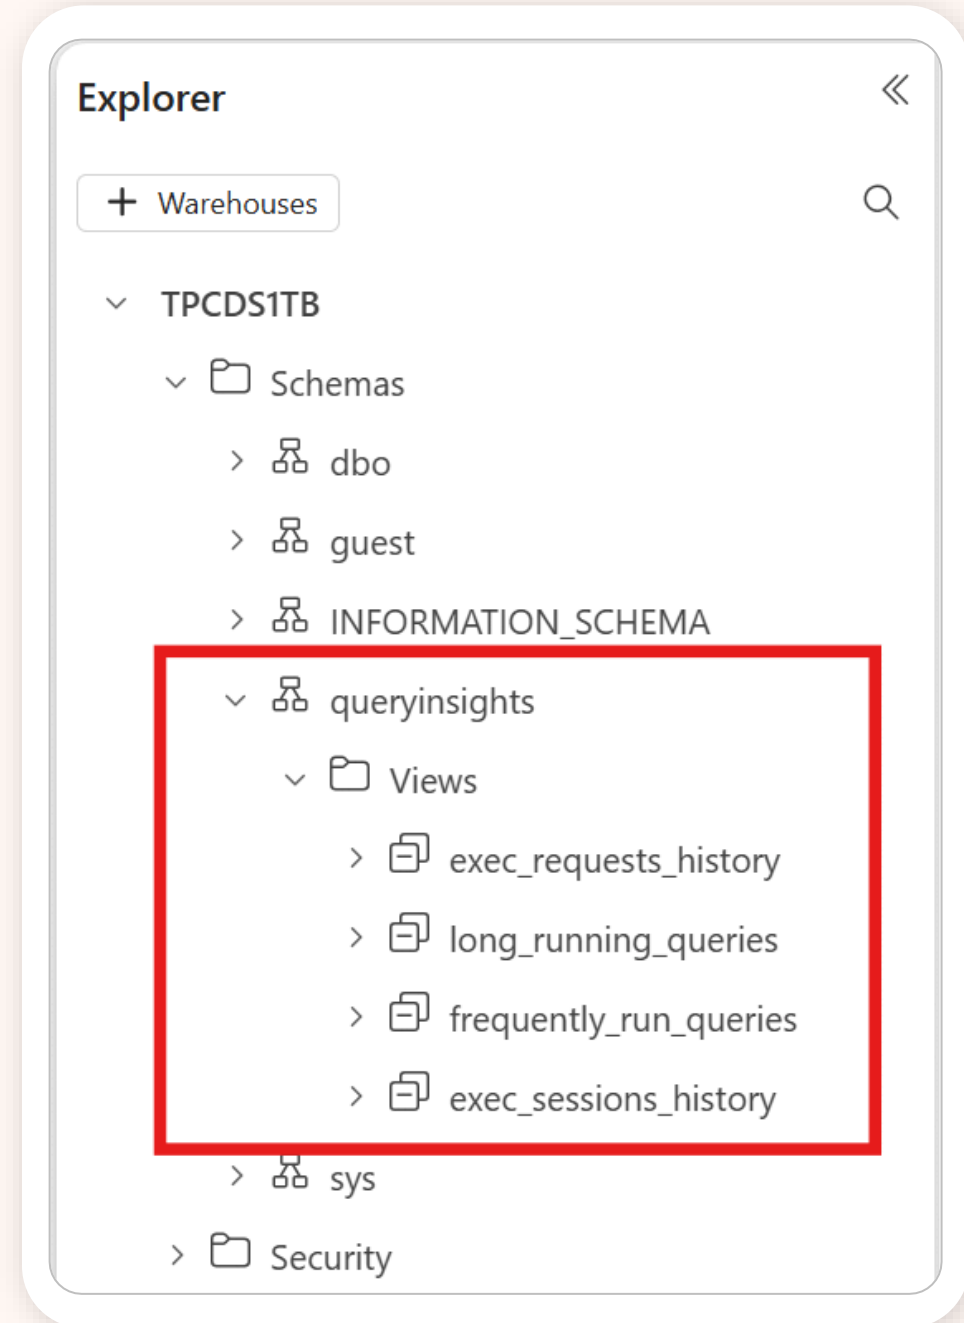
# Query Insights

Each query reflected in  
queryinsights.exec\_requests\_history

Populated on a delay of 5-15  
minutes.

Useful information about  
execution:

- Data scanned
- CPU time
- Historical runs
- Long running queries
- Highest cost queries



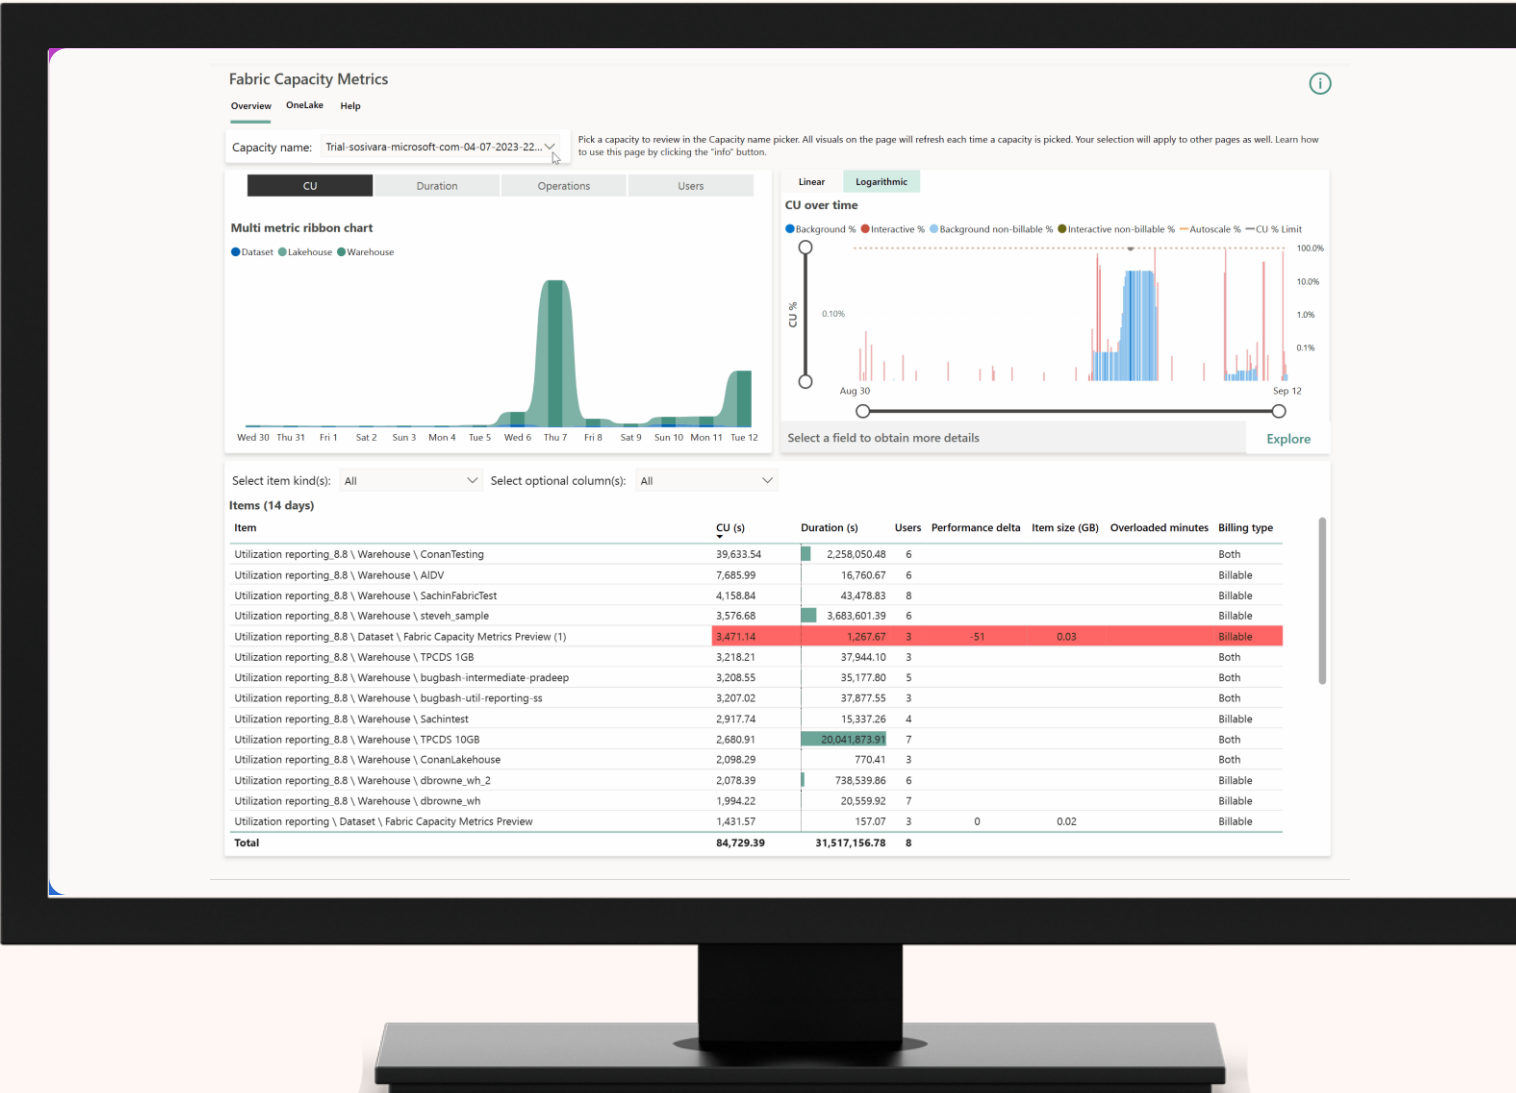
# Capacity Metrics

Visibility into capacity usage for all Fabric workloads, including data warehouse, in one place.

New usage is reflected in the report approximately every 10-15 minutes.

Use statement id to find individual T-SQL query.

The **Total CU (s)** field represents the total CU cost of the query.



# Identifying a query's cost

Optionally, with the total CU (s) of the query, compute the monetary cost of the query

Using the hourly rate from the Fabric pricing page, calculate the price per CU second for your capacity's region

Multiply the Total CU (s) value by the CU second cost to get the total query cost

Background operations

	Operation	Start	End	Status	User	Duration (s)	Total CU (s)	Timepoint CU (s)	Throttling (s)	% of Base Capacity	Billing type	Operation Id
\ Warehouse \ AIDV	Warehouse Query	9/7/2023 11:46:32 A...	9/7/2023 12:00:33 PM	InProgress	QueryC@microsoft.com	3819	5,528	1.92	0	0.10%	Billable	8D245338-6612
\ Warehouse \ CTesting	Warehouse Query	9/7/2023 9:45:26 AM	9/7/2023 10:02:33 AM	Success	QueryA@microsoft.com	938	5,323	1.85	0	0.10%	Billable	E21EFEEA-F638-
\ Warehouse \ CTesting	Warehouse Query	9/7/2023 10:44:43 A...	9/7/2023 10:59:49 AM	InProgress	QueryC@microsoft.com	819	5,252	1.82	0	0.09%	Billable	6898B839-EBC0-
\ Warehouse \ CTesting	Warehouse Query	9/7/2023 11:01:01 A...	9/7/2023 11:07:57 AM	Success	QueryA@microsoft.com	399	764	0.27	0	0.01%	Billable	C09DB8A9-382A-
\ Warehouse \ CTesting	Warehouse Query	9/7/2023 10:12:49 A...	9/7/2023 10:19:34 AM	Success	QueryC@microsoft.com	389	733	0.25	0	0.01%	Billable	445608F6-55DF-
\ Warehouse \ CTesting	Warehouse Query	9/7/2023 9:54:04 AM	9/7/2023 10:02:06 AM	Success	QueryA@microsoft.com	461	676	0.24	0	0.01%	Billable	C99021CC-42E2-
\ Warehouse \ CTesting	Warehouse Query	9/7/2023 10:16:48 A...	9/7/2023 10:20:49 AM	Success	QueryC@microsoft.com	163	228	0.08	0	0.00%	Billable	1D91D07D-CAC1-
\ Warehouse \ CTesting	Warehouse Query	9/7/2023 9:40:41 AM	9/7/2023 9:44:38 AM	InProgress	QueryA@microsoft.com	157	223	0.08	0	0.00%	Billable	7085EA77-AB7A-
\ Warehouse \ AIDV	Warehouse Query	9/7/2023 11:41:47 A...	9/7/2023 11:43:15 AM	InProgress	System	0	219	0.08	0	0.00%	Billable	55C346D-ED0C-
\ Warehouse \ CTesting	Warehouse Query	9/7/2023 10:37:57 A...	9/7/2023 10:39:22 AM	Success	QueryA@microsoft.com	70	126	0.04	0	0.00%	Billable	84D19CD9-F8A3-
\ Warehouse \ CTesting	Warehouse Query	9/7/2023 10:25:31 A...	9/7/2023 10:27:02 AM	InProgress	QueryC@microsoft.com	71	125	0.04	0	0.00%	Billable	A96439F6-C3FD-
\ Warehouse \ CTesting	Warehouse Query	9/7/2023 7:29:51 AM	9/7/2023 7:31:21 AM	Success	QueryA@microsoft.com	66	117	0.04	0	0.00%	Billable	BC8E1FAB-236D-
\ Warehouse \ AIDV	Warehouse Query	9/7/2023 11:44:25 A...	9/7/2023 12:01:18 PM	InProgress	QueryC@microsoft.com	934	111	0.04	0	0.00%	Billable	BD245338-6612
\ Warehouse \ CTesting	Warehouse Query	9/7/2023 10:27:48 A...	9/7/2023 10:29:03 AM	InProgress	QueryA@microsoft.com	55	106	0.04	0	0.00%	Billable	186731C0-9FF5-
\ Warehouse \ CTesting	Warehouse Query	9/7/2023 7:39:37 AM	9/7/2023 7:52:41 AM	Success	System	67	82	0.03	0	0.00%	Billable	3A14D71C-927B-
\ Warehouse \ CTesting	Warehouse Query	9/7/2023 7:39:37 AM	9/7/2023 8:10:58 AM	Success	System	60	73	0.03	0	0.00%	Billable	665AA23E-E3DC-
							6,546	27,853	12.08	0	0.63%	

## Cost Per Capacity Unit Second Calculation

- F2 Hourly Rate / 3,600 seconds per hour / 2 CU in an F2 = \$ per CU second
- F2 Hourly Rate in East US 2 = \$0.36
- $\$0.36 / 3600 / 2 = \$0.00005$  per CU second

## Total Query Cost Calculation

- \$ Per CU second \* Total CU (S) for the query = Total Query Cost
- $\$0.00005 * 5323 = \$0.26615$

# Monitoring Warehouse Activity

**What is happening right now?**

**DMVs**

**Fabric Data Warehouse  
Query Monitor**

**What happened in the past?**

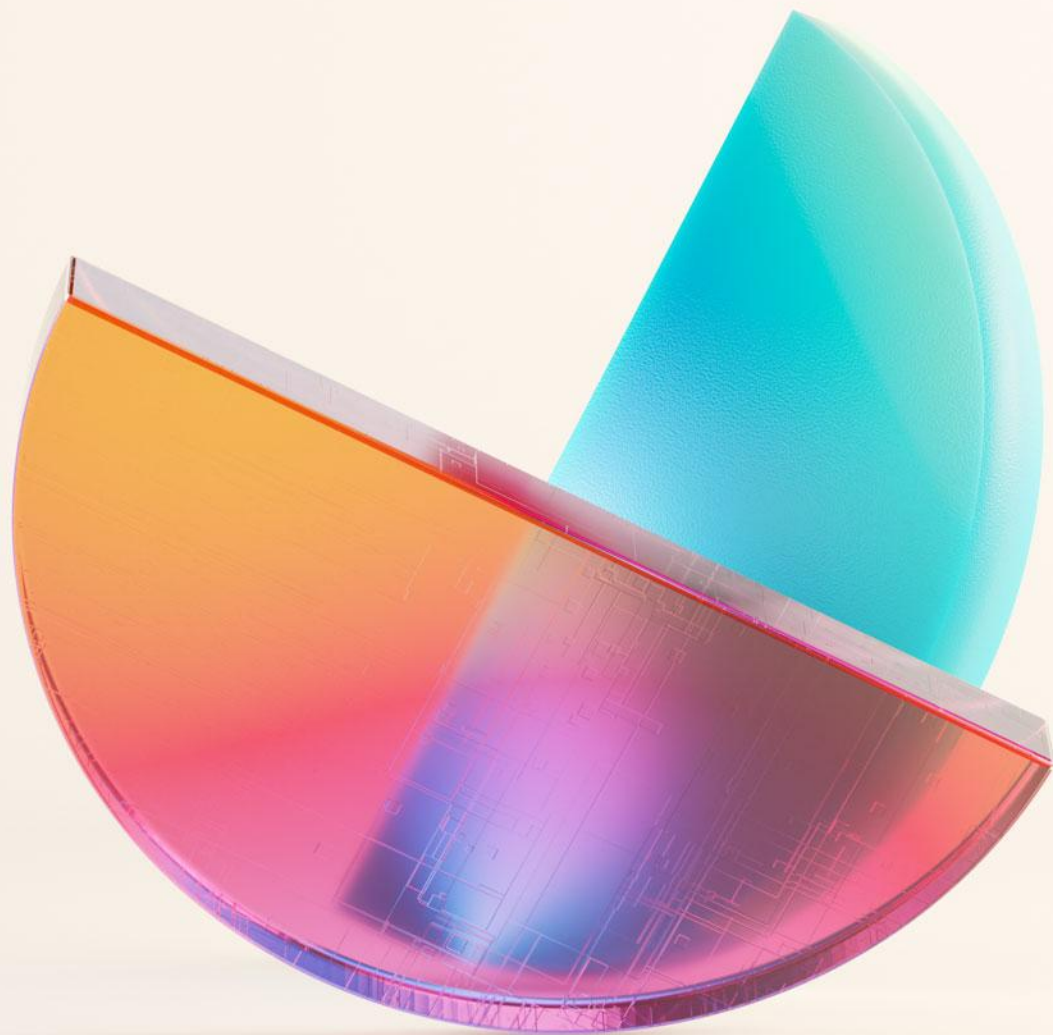
**Query Insights**

**Fabric Data Warehouse  
Query Monitor**

**What happened on my capacity?**

**Capacity Metrics**

**Query Insights\***

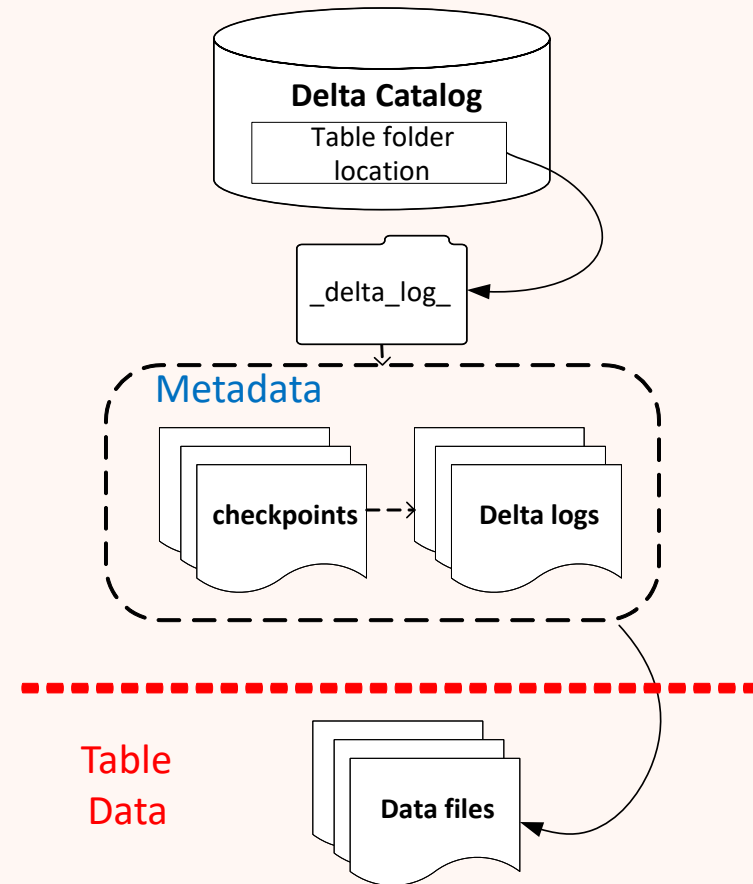


# Statistics and V-Order

## Module 7

# Storage engine design principles

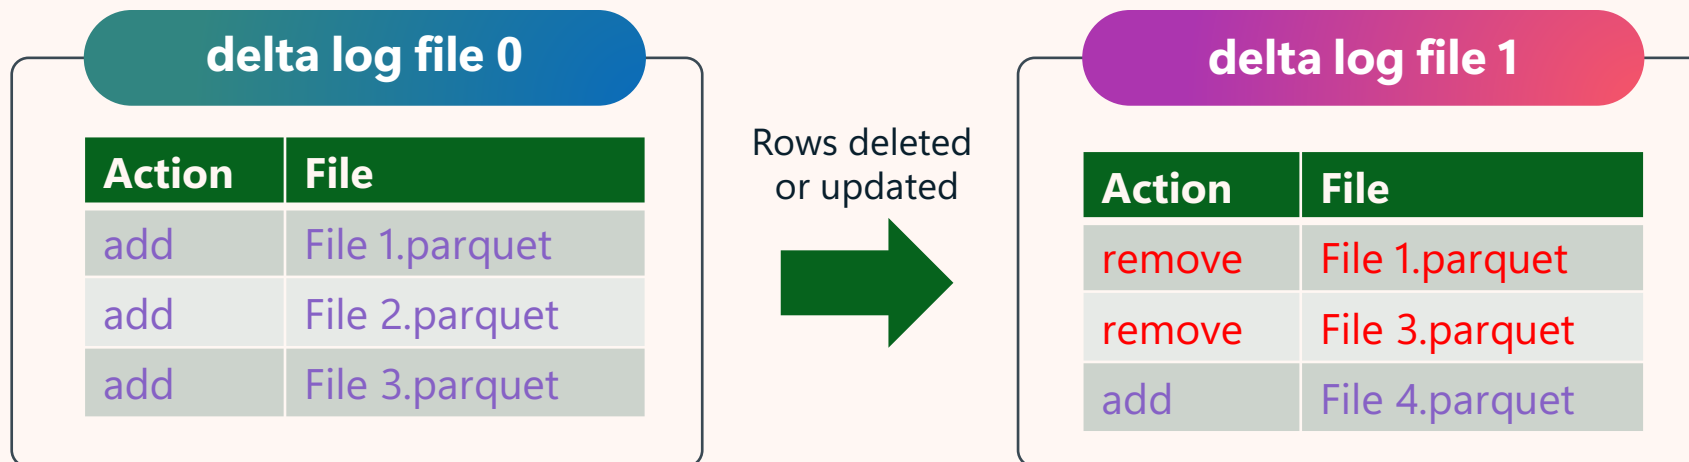
- Cloud first. Separation of State and Compute
- Columnar, immutable
- Open storage format
- Support for lineage-based features
- Snapshot Isolation semantics with multi-table transaction support
- No cross-component state sharing



# Data Compaction

As data changes, the warehouse engine maintains storage optimal

**No user action is required**



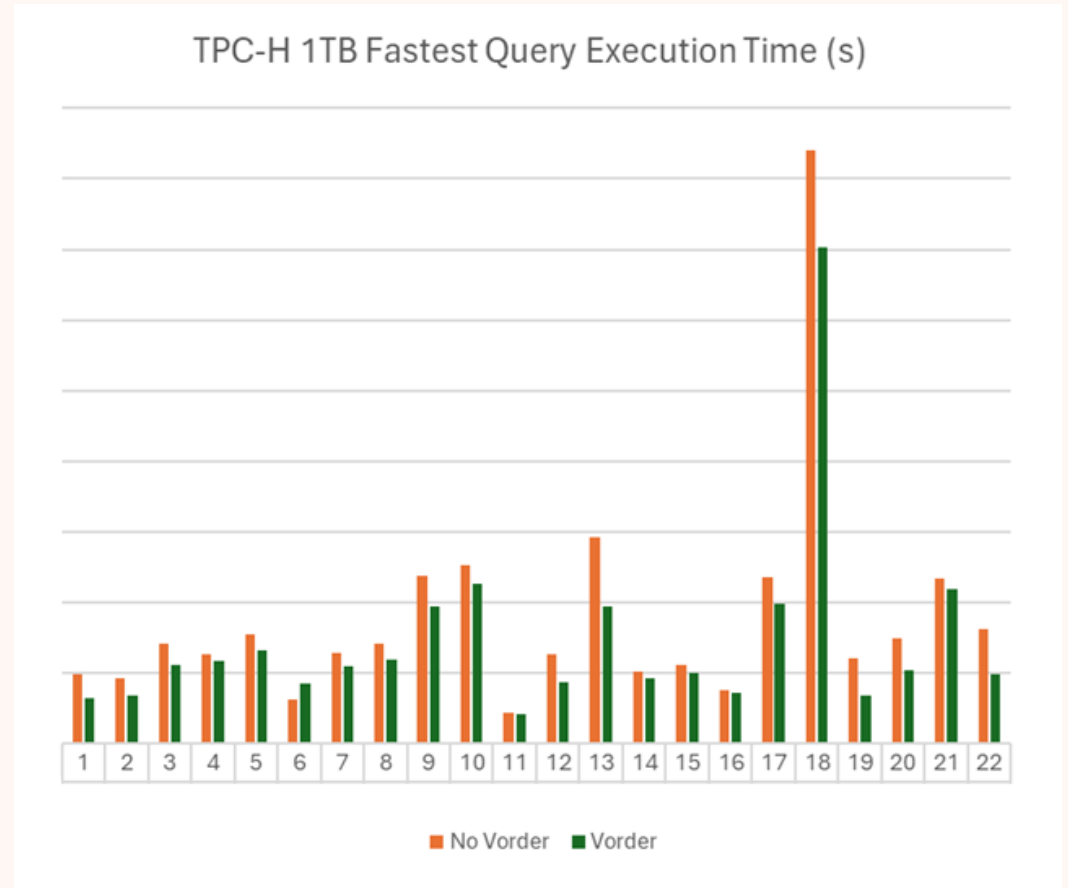


# V-Order

V-Order is a write time optimization to the parquet file format that enables lightning-fast reads under Microsoft Fabric compute engines, such as Power BI, SQL, Spark and others. It applies special sorting and compression to Parquet.

Warehouse queries benefit from faster read times with v-order, still ensuring files are 100% compliant to Parquet's open-source specification.

All warehouse data is written with v-order optimization during data ingestion.



Query performance: no V-Order vs V-Order

## V-Order (cont'd)

**IMPORTANT:** Disabling V-Order can only be done at the warehouse level, and it is irreversible. Once disabled, it cannot be enabled again. Users must consider all the performance impact of disabling V-Order before deciding to do so. **WE DO NOT RECOMMEND DISABLING V-ORDER WITHOUT A THOROUGH INVESTIGATION OF THE PERFORMANCE IMPACT.**

Disabling V-Order can make data ingestion faster, but it has the following impact:

- V-Order optimizes Parquet files for faster reads by applying sorting, row group distribution, dictionary encoding, and compression. Disabling it can lead to slower query performance, especially for analytics scenarios.
- Direct Lake mode in Power BI requires V-Order. If you disable V-Order, Power BI will fall back to direct query mode when querying tables in your V-Order-disabled warehouse.
- Only newly ingested data will not use V-Order; all existing files in your warehouse will remain with V-Order until rows are re-written (through ingestion or compaction).

# V-Order (cont'd)

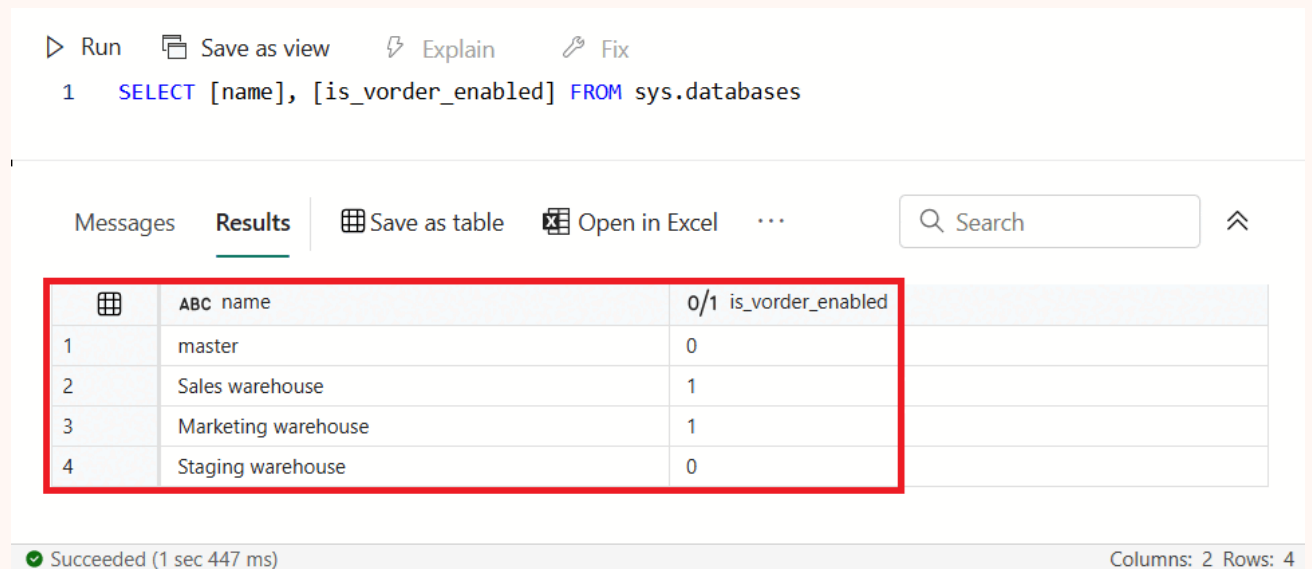
In special cases (e.g.: write intensive warehouses where reads are limited; staging warehouses; Power BI Direct Lake mode not needed), V-ORDER can be disabled.

To disable V-Order in the current Warehouse:

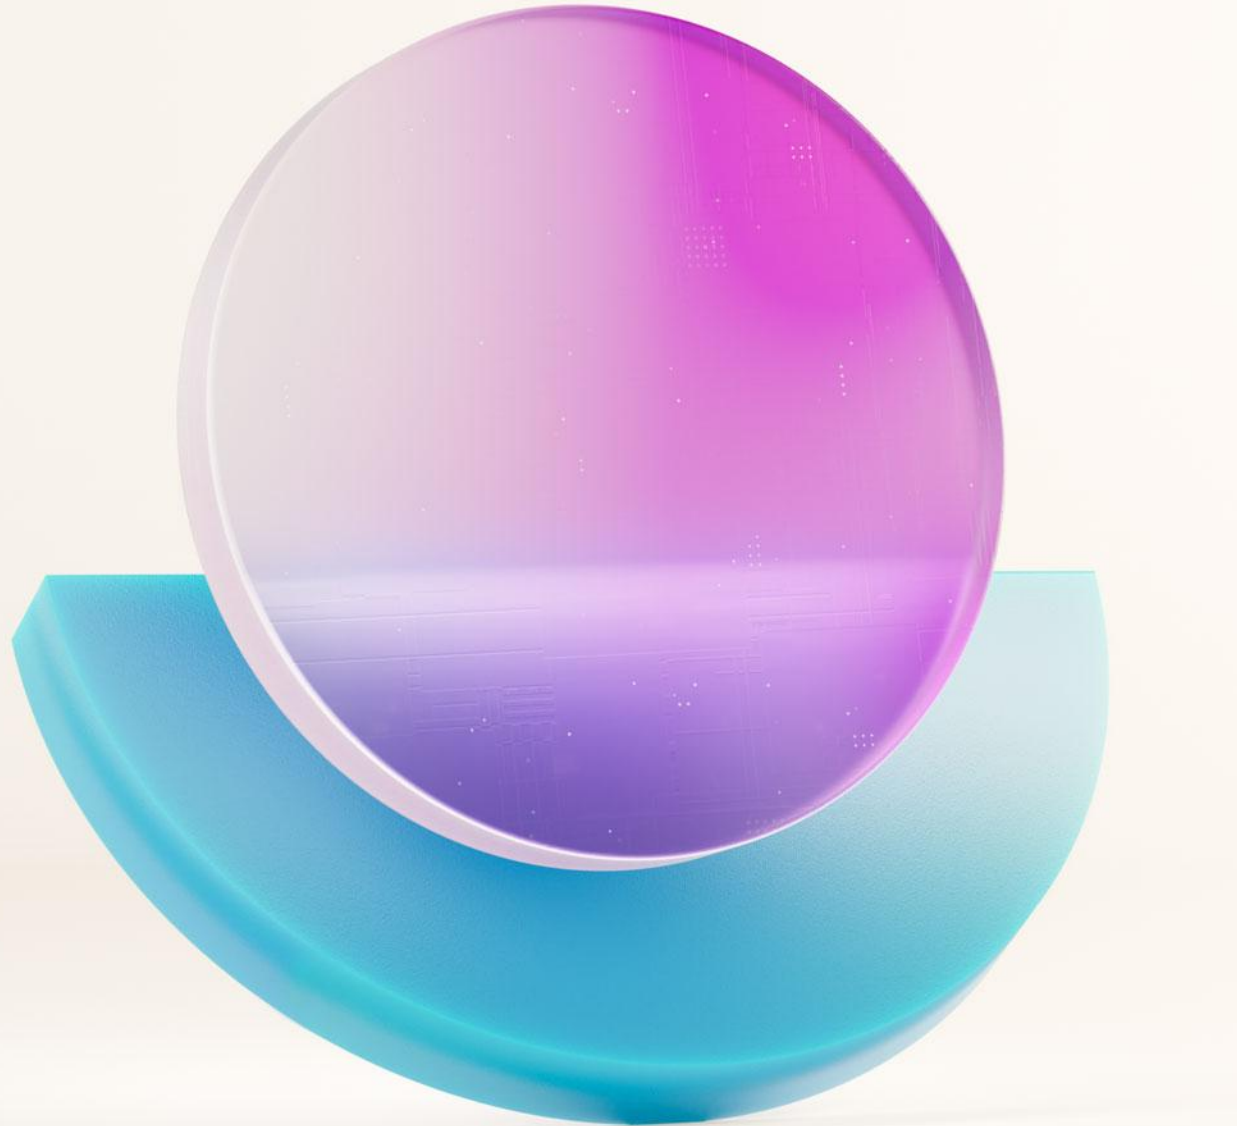
```
ALTER DATABASE CURRENT SET VORDER = OFF
```

To check the current configuration of V-ORDER on your warehouse:

```
SELECT [name], [is_vorder_enabled]  
FROM sys.databases
```



	name	is_vorder_enabled
1	master	0
2	Sales warehouse	1
3	Marketing warehouse	1
4	Staging warehouse	0



# Optimizing queries

# Managing Statistics

Statistics are objects that contain relevant information about your data. They enable the query optimizer to estimate a query plan with the least system overhead (I/O, CPU) possible, resulting in faster query execution.

## **Automatic statistics - RECOMMENDED**

- Engine automatically creates and maintains statistics at query time
- Because these automatic operations are done synchronously, expect the query duration to include this time if the needed statistics do not yet exist, or significant data changes have happened since the last statistics refresh

## **User-defined statistics**

- The user manually uses data definition language (DDL) syntax to create, update, and drop statistics as needed
- Consider updating column-level statistics regularly after data updates that significantly change the row count or distribution of the data

# Manually managing statistics

If creating statistics manually, consider focusing on columns heavily used in your query workload (specifically in GROUP BYs, ORDER BYs, filters, and JOINS). Examples:

Create statistics for the CustomerKey column of a table:

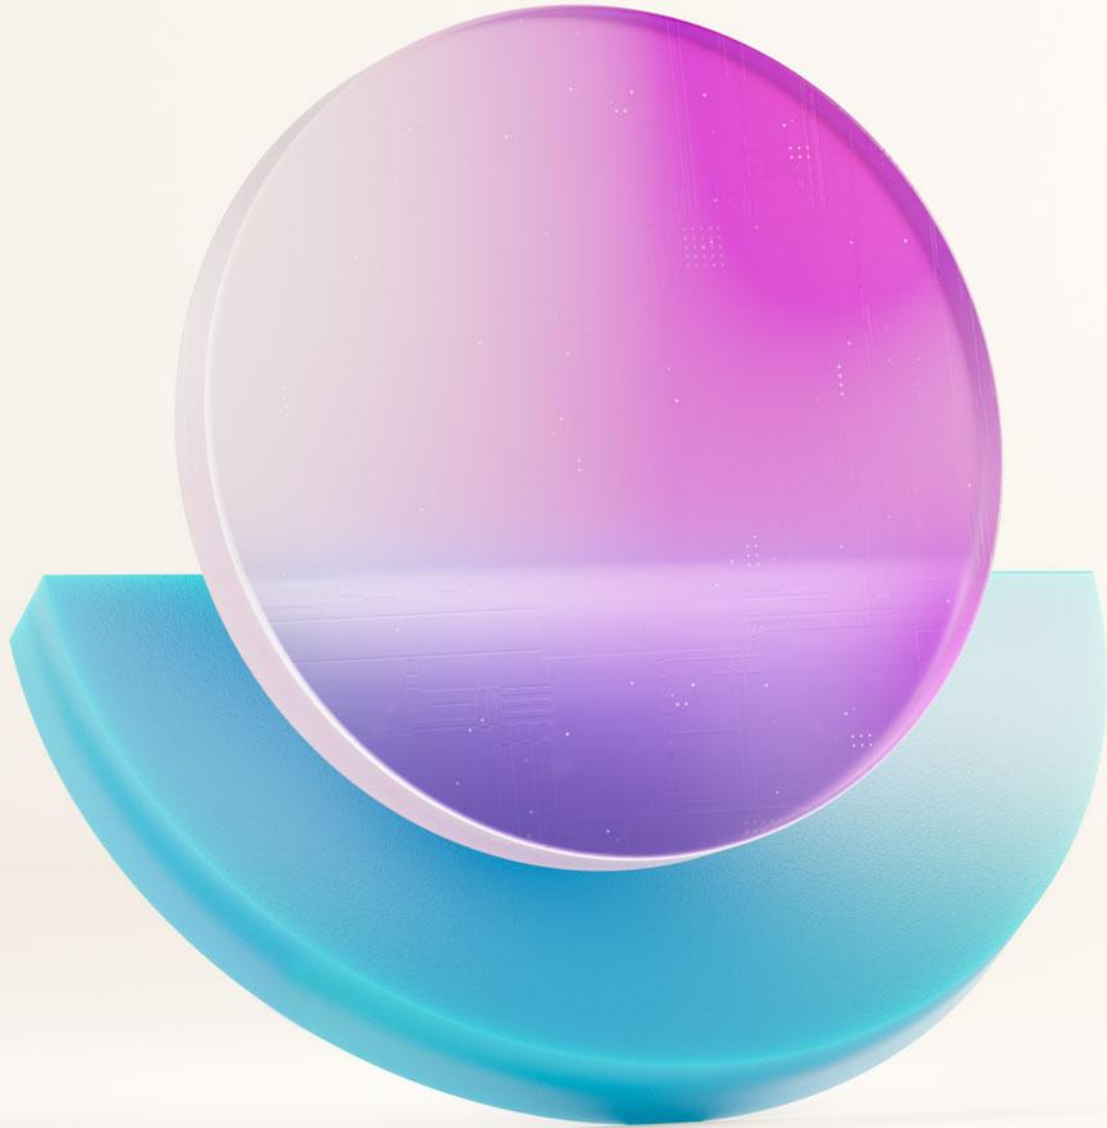
```
CREATE STATISTICS DimCustomer_CustomerKey_FullScan  
ON dbo.DimCustomer (CustomerKey) WITH FULLSCAN;
```

Update statistics for the statistics object previously created (perhaps after a large data update):

```
UPDATE STATISTICS dbo.DimCustomer (DimCustomer_CustomerKey_FullScan) WITH FULLSCAN;
```

To show information about the histogram of the statistics object:

```
DBCC SHOW_STATISTICS ("dbo.DimCustomer", "DimCustomer_CustomerKey_FullScan") WITH HISTOGRAM;
```



Demo

# Obtaining the query plan (preview)

Now in preview: `SET SHOWPLAN_XML { ON | OFF }`

- Using SET SHOWPLAN\_XML helps you analyze and optimize complex queries without executing them, ensuring you can make informed decisions about query tuning.
- Fabric Data Warehouse returns detailed information about with the query plan for query execution.

The screenshot displays the Microsoft Fabric Data Warehouse interface. At the top, there are buttons for 'Run', 'Save as view', 'Explain query', and 'Fix query errors'. Below these is a 'Preview' banner with a warning: 'Copilot uses AI. Mistakes can happen. Verify code suggestions before running them. [Review terms](#)'. The main area contains a SQL query:

```
1 SET SHOWPLAN_XML ON;  
2 GO  
3  
4 SELECT O_CUSTKEY, O_CLERK, COUNT(*) as ORDER_COUNT  
5 FROM Orders  
6 WHERE O_CUSTKEY > 1106283 AND O_CUSTKEY < 21062833124  
7 GROUP BY O_CUSTKEY, O_CLERK
```

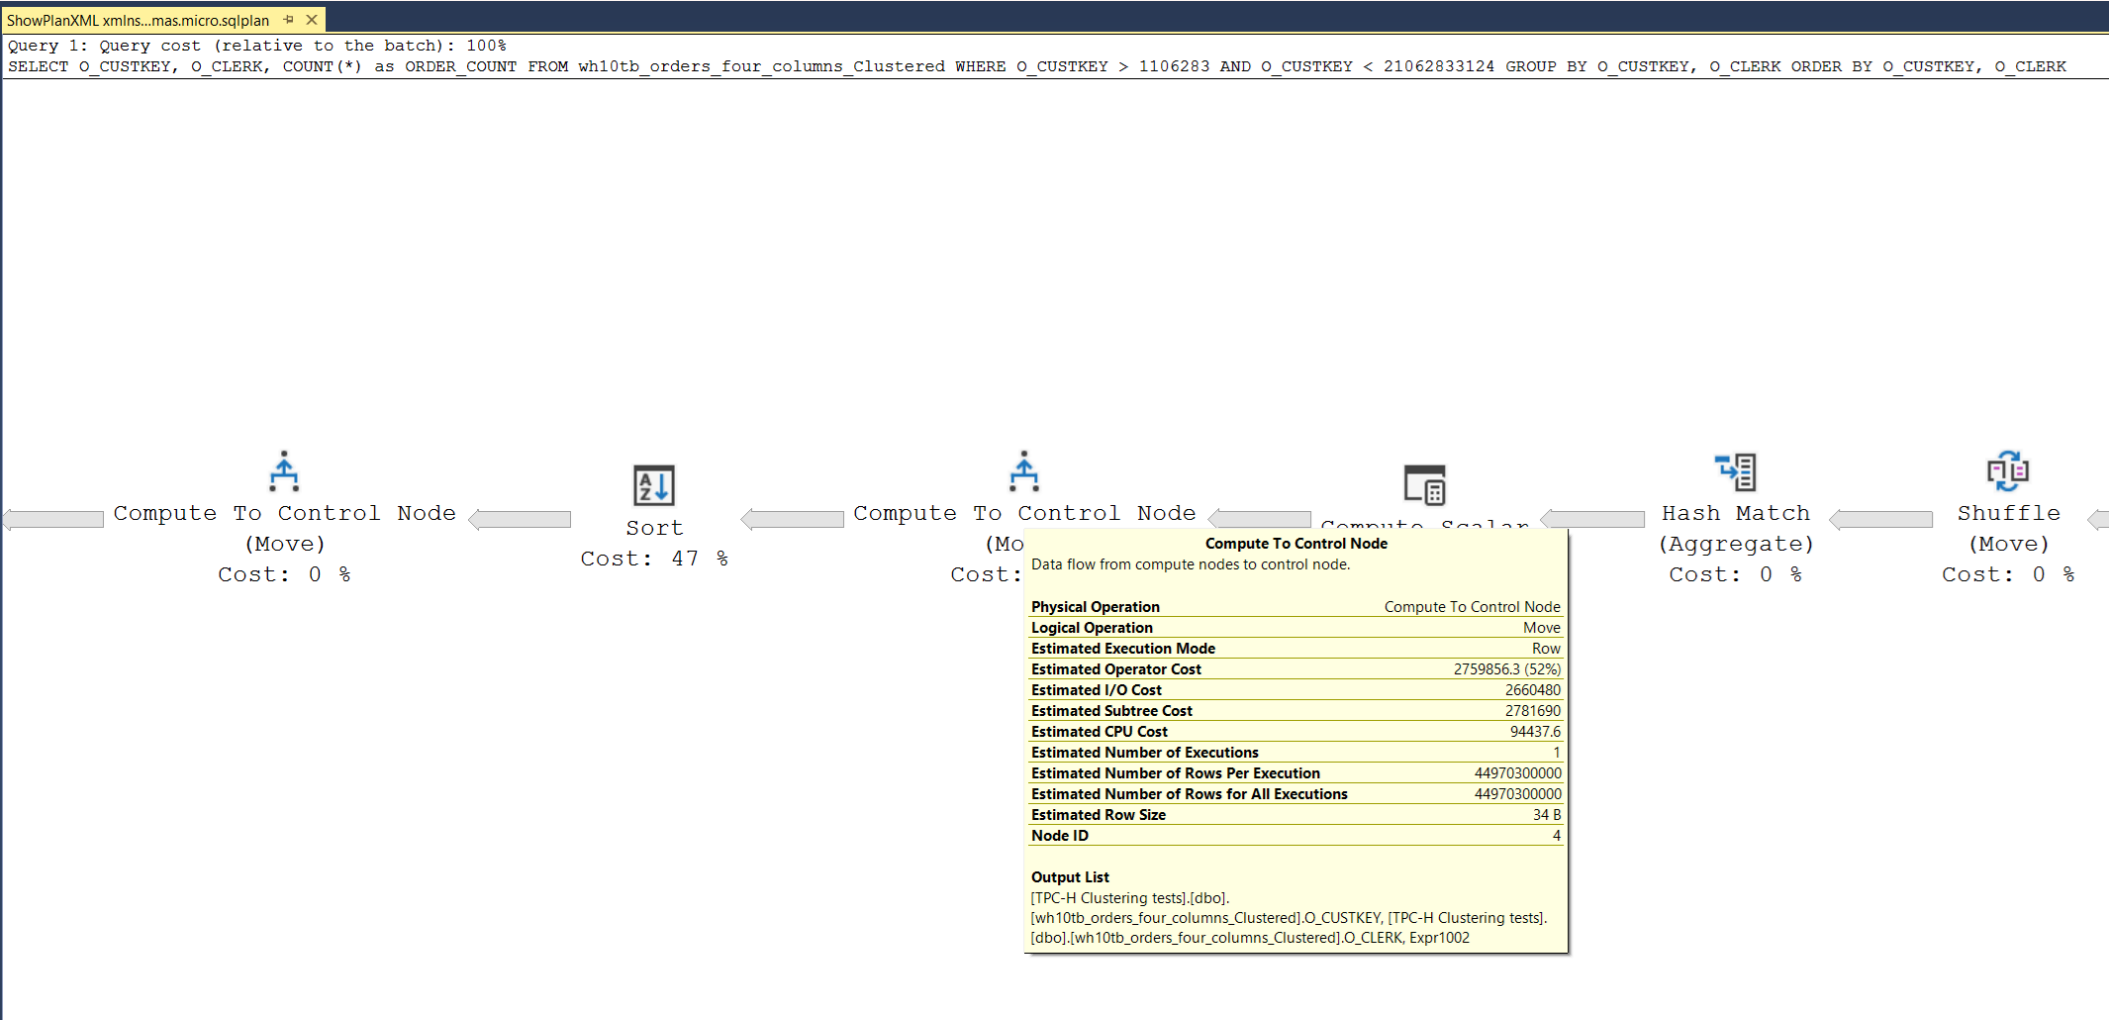
Below the query, there are tabs for 'Messages' and 'Results'. The 'Results' tab is active, showing a table with one row of XML data. The XML is a Microsoft SQL Server 2005 XML Showplan, starting with: `<ShowPlanXML xmlns="http://schemas.microsoft.com/sqlserver/2004/07/showplan" Version="1.595" Build="17.0.509.106"><BatchSequence><Batch><Statements><StmtSimple StatementText="&#xa...`

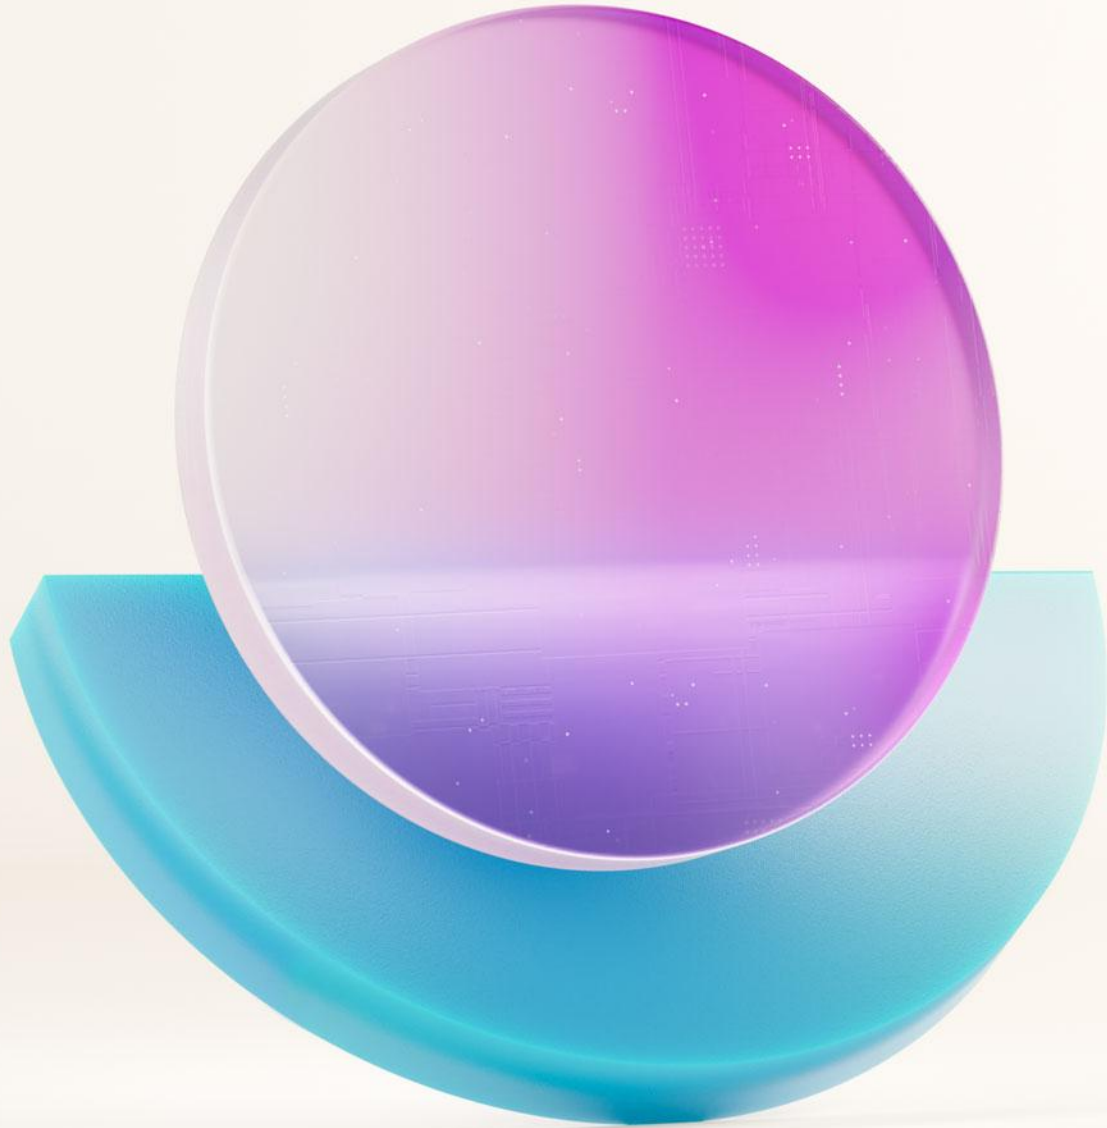
At the bottom, a status bar shows 'Succeeded (0 sec 969 ms)' and 'Copilot completions: On' and 'Copilot completions: Ready'. The bottom right corner indicates 'Columns: 1 Rows: 1'.



# Obtaining the query plan (preview) (cont'd)

Query plans can be visualized in SQL Server Management Studio

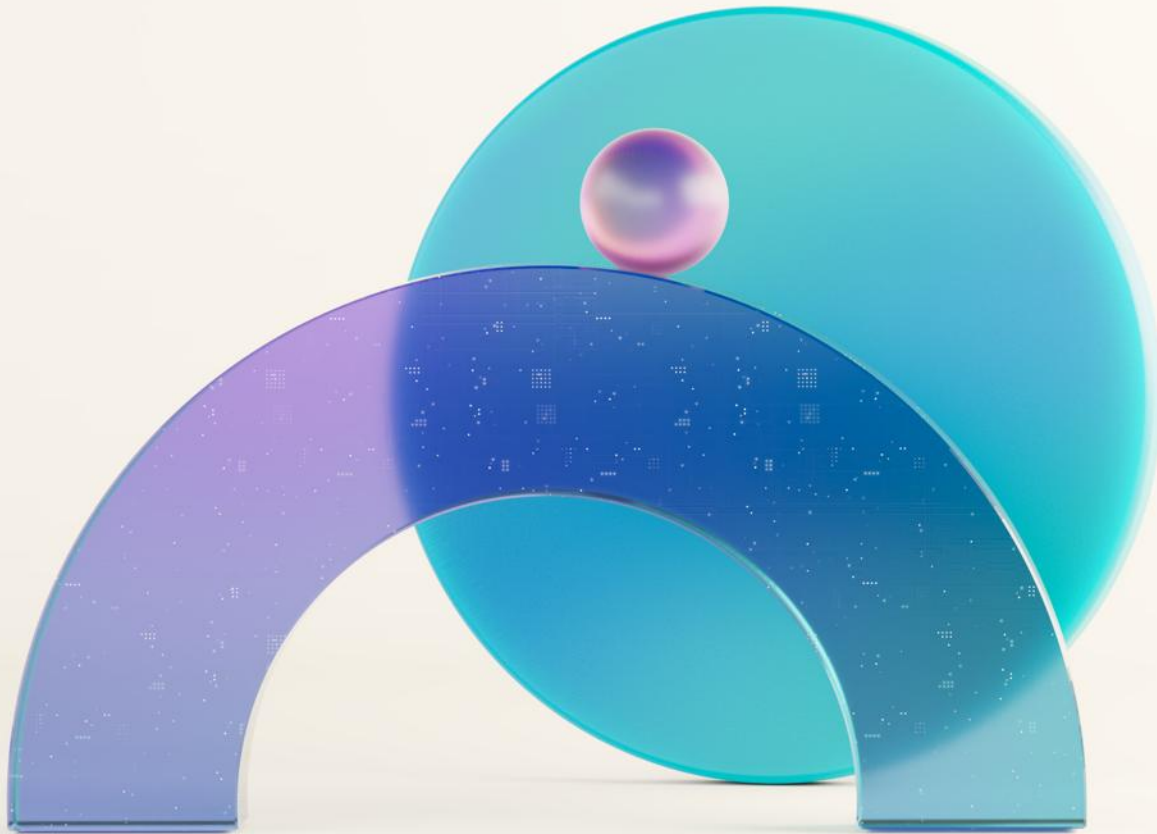




Demo



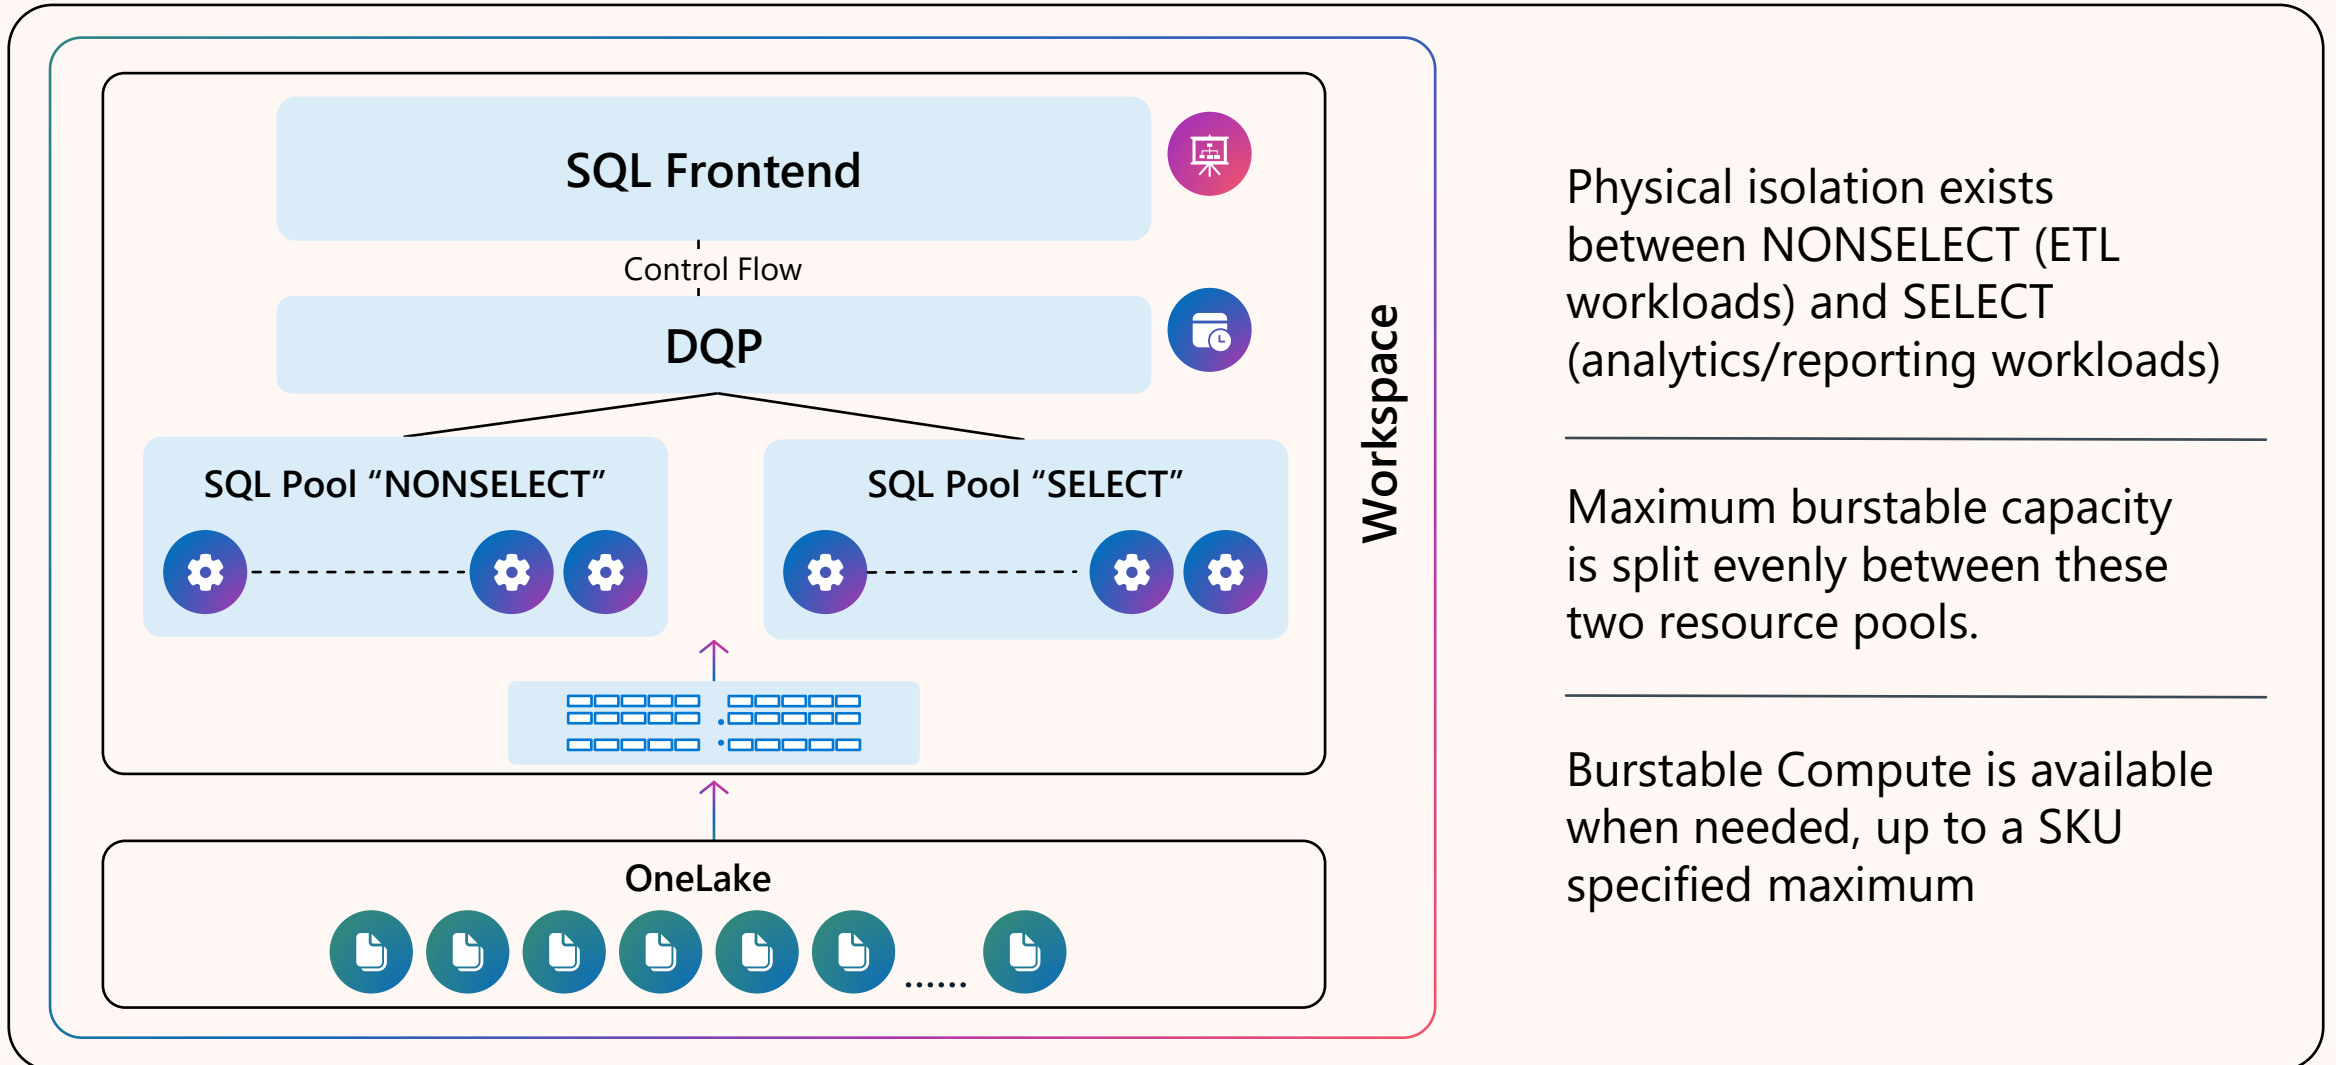
10-minute break



# Workload Management

## Module 8

# Intelligent Workload Management

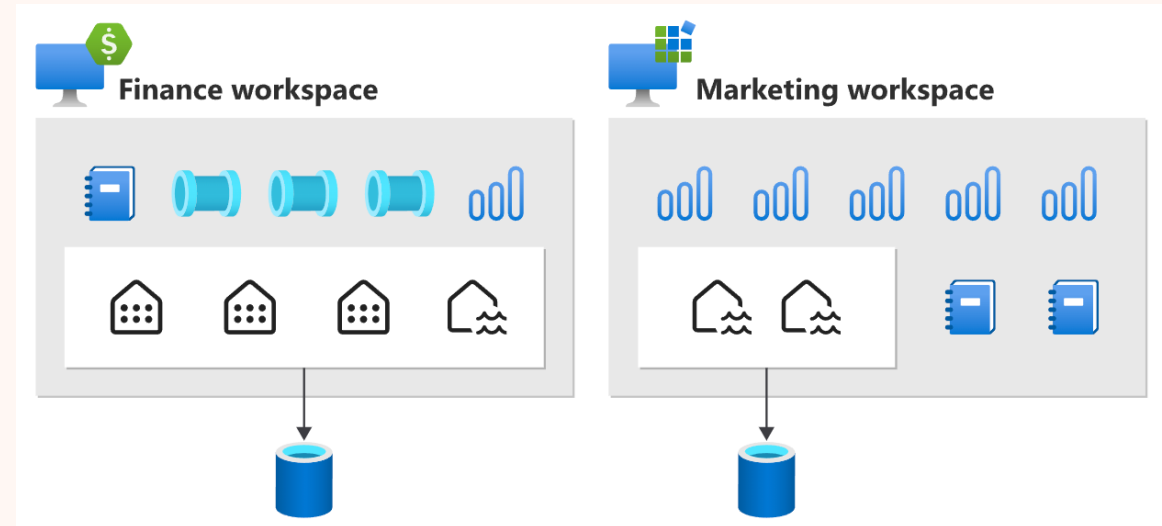


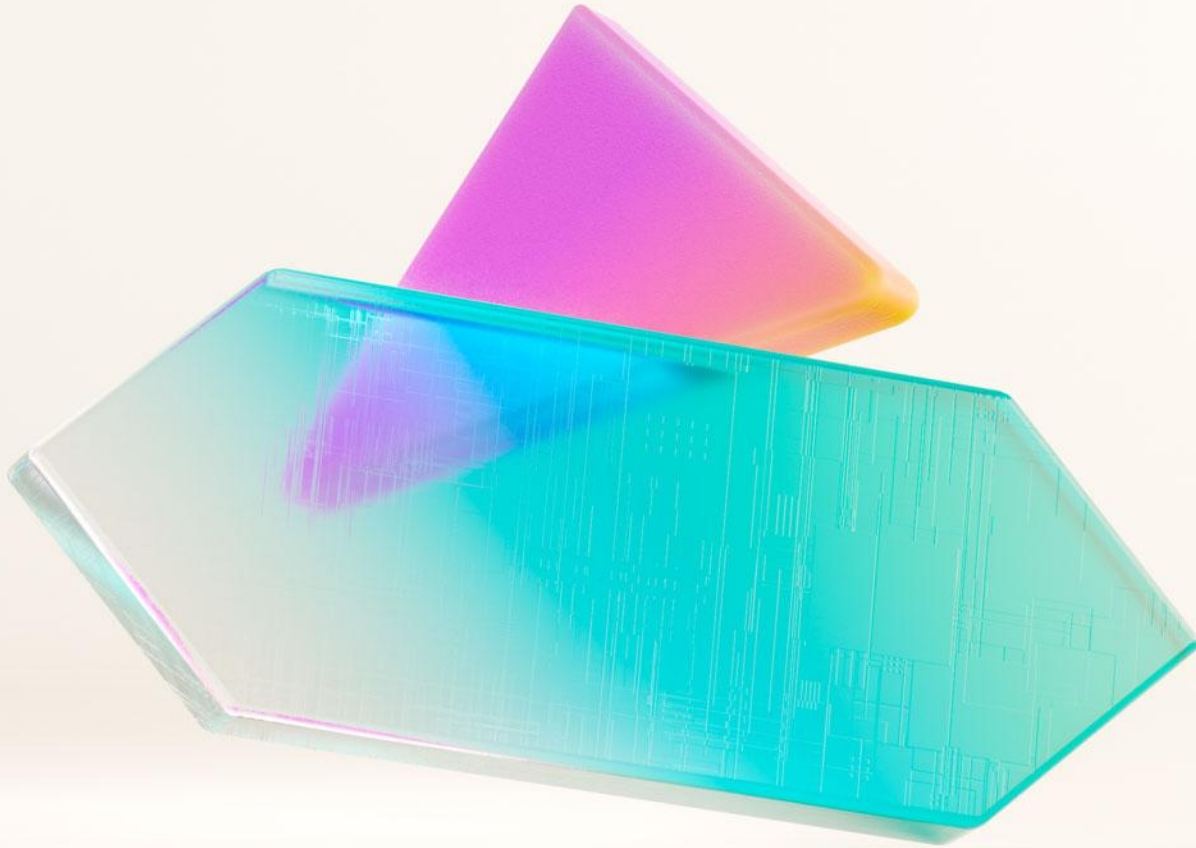
# Workspace isolation for different workloads

Microsoft Fabric workspaces provides a natural isolation boundary of the distributed compute system. Workloads can take advantage of this boundary to manage both cost and performance.

OneLake shortcuts can be used to create read-only replicas of tables in other workspaces to distribute load across multiple SQL engines, creating an isolation boundary.

This can effectively increase the maximum number of sessions performing read-only queries.





# Security

## Module 9

# Permission model

Microsoft Fabric permissions and granular SQL permissions work together to govern Warehouse access and the user permissions once connected

Workspace roles provide Microsoft Fabric permissions for all warehouses within a workspace

Fabric capability	Admin	Member	Contributor	Viewer
Delete the workspace	✓	✗	✗	✗
Add admins	✓	✗	✗	✗
Add members	✓	✓	✗	✗
Write data	✓	✓	✓	✗
Create items	✓	✓	✓	✗
Read data	✓	✓	✓	✓

Learn more about workspace roles on <https://learn.microsoft.com/en-us/fabric/fundamentals/roles-workspaces>



# Workspace roles in Fabric Data Warehouse

Assigning users to the various workspace roles uses the following SQL permissions

Workspace role	Description
<b>Admin</b>	<p>Grants the user <b>CONTROL</b> access for each Warehouse and SQL analytics endpoint within the workspace, providing them with full read/write permissions and the ability to manage granular user SQL permissions.</p> <p>Allows the user to see workspace-scoped session, monitor connections and requests in DMVs via TSQL, and KILL sessions.</p>
<b>Member</b>	<p>Grants the user <b>CONTROL</b> access for each Warehouse and SQL analytics endpoint within the workspace, providing them with full read/write permissions and the ability to manage granular user SQL permissions. Share all content within the warehouse.</p>
<b>Contributor</b>	<p>Grants the user <b>CONTROL</b> access for each Warehouse and SQL analytics endpoint within the workspace, providing them with full read/write permissions and the ability to manage granular user SQL permissions.</p>
<b>Viewer</b>	<p>Grants the user <b>CONNECT</b> and ReadData permissions for each Warehouse and SQL analytics endpoint within the workspace. Viewers have SQL permissions to read data from tables/views using T-SQL.</p>

# Granular security on SQL objects

- Object-level-security can be managed using GRANT, REVOKE, and DENY T-SQL syntax
- Users can be assigned to SQL roles, both custom and built-in database roles

To grant users access to database objects:


```
GRANT EXECUTE ON OBJECT::dbo.UpdateDimCity TO [roger@contoso.com]
```

Seamlessly, to deny access:

```
DENY EXECUTE ON OBJECT::dbo.UpdateDimCity TO [roger@contoso.com]
```

# Row-level security

- Enables you to use group membership or execution context to control access to rows in a database table.
- Examples:
  - Ensure that workers access only those data rows that are pertinent to their department
  - Allow sales managers access to a global sales table, but only to rows pertinent to their region

	SaleID	CustomerID	TotalAmmount	Date	Region
	1001	3345	121.00	3/29/2025	Southeast
	1002	2561	629.00	3/29/2025	Southeast
					
					

# Row-level security

-- Creating a function for the SalesRep evaluation

```
CREATE FUNCTION Security.tvf_securitypredicate(@SalesRep AS nvarchar(50))  
    RETURNS TABLE  
WITH SCHEMABINDING  
AS  
    RETURN SELECT 1 AS tvf_securitypredicate_result  
WHERE @SalesRep = USER_NAME() OR USER_NAME() = 'manager@contoso.com';
```

-- Using the function to create a Security Policy

```
CREATE SECURITY POLICY SalesFilter  
ADD FILTER PREDICATE Security.tvf_securitypredicate(SalesRep)  
ON sales.Orders  
WITH (STATE = ON);
```

# Column-level security

- Enables you to use group membership or execution context to control access to columns in a database table.
- Examples:
  - Ensure that workers can't see sensitive information about a customer (e.g.: SSN, credit card number)
  - Allow HR personnel to see employee names, but hide salaries



First name	Last name	Job role	Manager ID	Salary
Dany	Bélisle	Data Engineer	627819	
Camelia	Banica	Data Analyst	726541	
Emil	Moldovan	SQL Developer	726541	
Ioana	Costache	Data guru	019275	

# Column-level security

```
GRANT SELECT ON Employees(FirstName, LastName, JobRole, ManagerID) TO  
[charlie@contoso.com]
```

Attempts to read the Salary column will fail:

```
SELECT * FROM Employees
```

Msg 230, Level 14, State 1, Line 12

The SELECT permission was denied on the column 'Salary' of the object 'Employees', database 'HumanResources', schema 'dbo'.

# Dynamic data masking

- Limits sensitive data exposure by masking it to nonprivileged users. It can be used to greatly simplify the design and coding of security in your application
- The data isn't changed, so it can be used with existing applications since masking rules are applied to query results
- Minimal (if any) effect in the application layer

First name	Last name	E-mail	Phone	Credit Card Number
Dany	Bélisle	<a href="mailto:dXXXXXX@XXXXXX.com">dXXXXXX@XXXXXX.com</a>	(425) ###-1254	XXXX-XXXX-XXXX-1234
Camelia	Banica	<a href="mailto:cXXXXXX@XXXXXX.com">cXXXXXX@XXXXXX.com</a>	(425) ###-7284	XXXX-XXXX-XXXX-4321
Emil	Moldovan	<a href="mailto:eXXXXXX@XXXXXX.com">eXXXXXX@XXXXXX.com</a>	(425) ###-8271	XXXX-XXXX-XXXX-0123
Ioana	Costache	<a href="mailto:IXXXXXX@XXXXXX.com">IXXXXXX@XXXXXX.com</a>	(425) ###-0021	XXXX-XXXX-XXXX-3210

# Dynamic data masking

```
CREATE TABLE dbo.EmployeeData (  
    EmployeeID INT  
    ,FirstName VARCHAR(50) MASKED WITH (FUNCTION = 'partial(1,"-",2)') NULL  
    ,LastName VARCHAR(50) MASKED WITH (FUNCTION = 'default()') NULL  
    ,SSN CHAR(11) MASKED WITH (FUNCTION = 'partial(0,"XX-XX-",4)') NULL  
    ,email VARCHAR(256) NULL);
```

In this example:

- The **FirstName** column shows only the first and last two characters of the string, with - in the middle.
- The **LastName** column shows **XXXX**.
- The **SSN** column shows **XXX-XX-** followed by the last four characters of the string.